



Operating Systems (A)

(Honor Track)

Lecture 13: Scheduling (II)

Yao Guo (郭耀)

Peking University

Fall 2021



Review: Scheduling

Scheduling objectives:

- Performance
- Fair
- Priority
- Encourage good behavior
- Support heavy loads
- Adapt to different environments
 - interactive, real-time, multi-media



Single Processor Scheduling

- Scheduling for batch systems
 - First Come First Serve (FCFS)
 - Shortest Job First (SJF)
- Interactive scheduling
 - Round Robin (RR)
 - Priority Scheduling
 - Priority inversion
 - Priority inheritance
 - Multiple-Level Feedback Queues (MLFQ)



This Lecture

Multiprocessor and Real-Time Scheduling

Multiprocessor scheduling
Real-time scheduling



Buzz Words

Load sharing

Gang scheduling

Hard/soft real-time task

Periodic/aperiodic task

Deadline scheduling

**Rate monotonic
scheduling (RMS)**



This Lecture

Multiprocessor and Real-Time Scheduling

Multiprocessor scheduling

Real-time scheduling



Classifications of Multiprocessor Systems

- Loosely coupled processors
 - Each has their memory & I/O channels
- Functionally specialized processors
 - Such as I/O processor
 - Controlled by a master processor
- Tightly coupled multiprocessor
 - Processors share main memory
 - Controlled by operating system
- We will mainly talk about tightly coupled multiprocessor in this lecture



Assignment of Processes to Processors

- The simplest way: treat processors as a pooled resource and assign processes to processors on demand
- Should the assignment be static or dynamic?
- Dynamic assignment
 - Threads can move from a queue for one processor to a queue for another processor
 - E.g., Linux



Static Assignment

- Permanently assign processes to a processor
 - Dedicated short-term queue for each processor
 - Less overhead
 - Allows the use of 'group' or 'gang' scheduling (see later)
- But may leave a processor idle, while others have a backlog
- Solution: use a **common queue**
 - All processes go into one global queue and are scheduled to any available processor



Assignment of Processes to Processors

- Both dynamic and static methods require some way of assigning a process to a processor
- Two methods:
 - Master/Slave
 - Peer
- There are of course a spectrum of approaches between these two extremes



Master / Slave Architecture

- Key kernel functions always run on a particular processor
- **Master**: responsible for scheduling
- **Slave**: sends service request to the master
- Disadvantages
 - Failure of master brings down whole system
 - Master can become a performance bottleneck

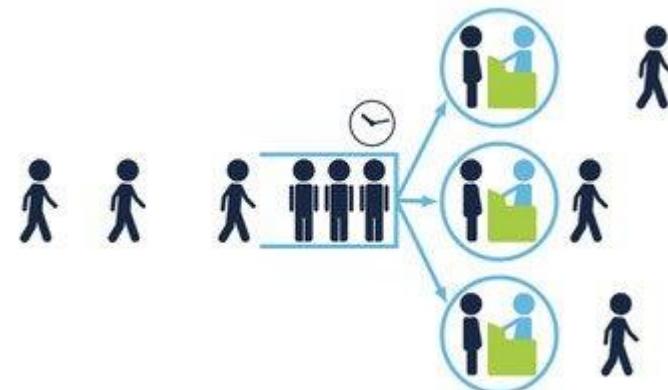


Peer Architecture

- Kernel can execute on any processor
- Each processor does **self-scheduling** from the pool of available processes.
- **Complicates the operating system**
 - The operating system must ensure that two processors do not choose the same process and that the processes are not somehow lost from the queue.
 - Techniques must be employed to resolve and synchronize competing claims to resources.

Process Scheduling

- Usually processes are not dedicated to processors
- A single queue is used for all processes
- Or use multiple queues based on priority
 - All queues feed to the common pool of processors
- We can view the system as being a multiserver queuing architecture.





Thread Scheduling

- Threads execute separately from the rest of the process
 - An application can be implemented as a set of threads that cooperate and execute concurrently in the same address space
- Threads can be used to exploit true parallelism in an application
 - Dramatic gains in performance are possible in multi-processor systems, compared to running in uniprocessor systems



Approaches to Thread Scheduling

- Many proposals exist but four general approaches stand out:
 - Load Sharing
 - Gang Scheduling
 - Dedicated processor assignment
 - Dynamic scheduling



Load Sharing

- A global queue of ready threads is maintained
 - In the (shared) main memory
- Each processor, when idle, selects a thread from the queue

- Load is distributed evenly across the processors
- No centralized scheduler required



Disadvantages of Load Sharing

- The central queue needs **mutual exclusion**
 - It may become a bottleneck if many processors look for work at the same time.
- Preempted threads are unlikely to resume execution on the same processor
 - Cache issues...
- If all threads are in the global queue, **all threads of a program** will not be likely to gain access to the processors at the same time



Gang Scheduling

- A set of related threads is scheduled to run on a set of processors at the same time
 - If closely related processes execute in parallel, synchronization blocking may be reduced, **less process switching** may be necessary, and performance will increase.
 - **Scheduling overhead may be reduced** because a single decision affects a number of processors and processes at one time.

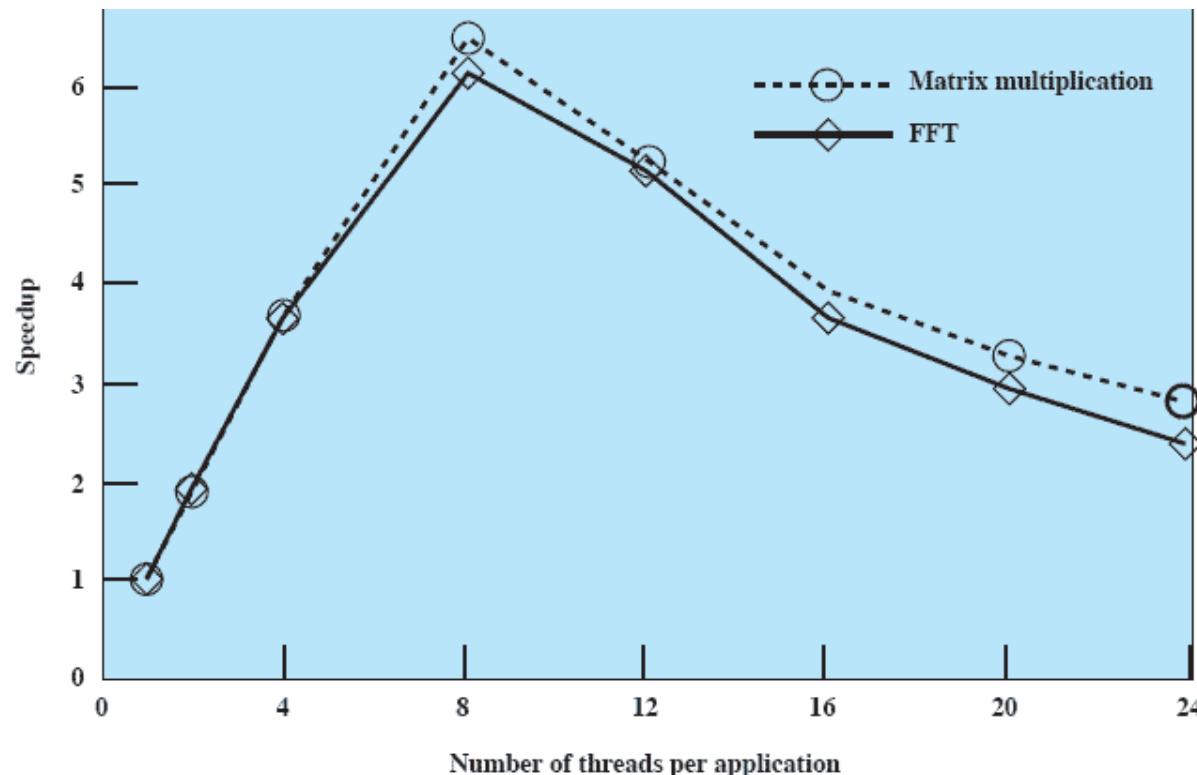


Dedicated Processor Assignment

- When an application is scheduled, each of its threads is assigned to a processor
- Some processors may be idle
 - No multiprogramming of processors (time-multiplexing)
- ***But***
 - In *highly parallel systems*, processor utilization is less important than effectiveness
 - Avoiding process switching speeds up programs

Dynamic Scheduling

- The number of threads in a process are altered dynamically by the application
 - This allows OS to adjust the load to improve utilization





This Lecture

Multiprocessor and Real-Time Scheduling

Multiprocessor scheduling

Real-time scheduling



Real-Time Scheduling

- Real-time computing:
 - Correctness of the system depends not only on the logical result of the computation but also on the time at which the results are produced
- Tasks or processes attempt to control or react to events that take place in the outside world
- These events occur in “real-time” and tasks must be able to keep up with them



Hard vs Soft Real-time

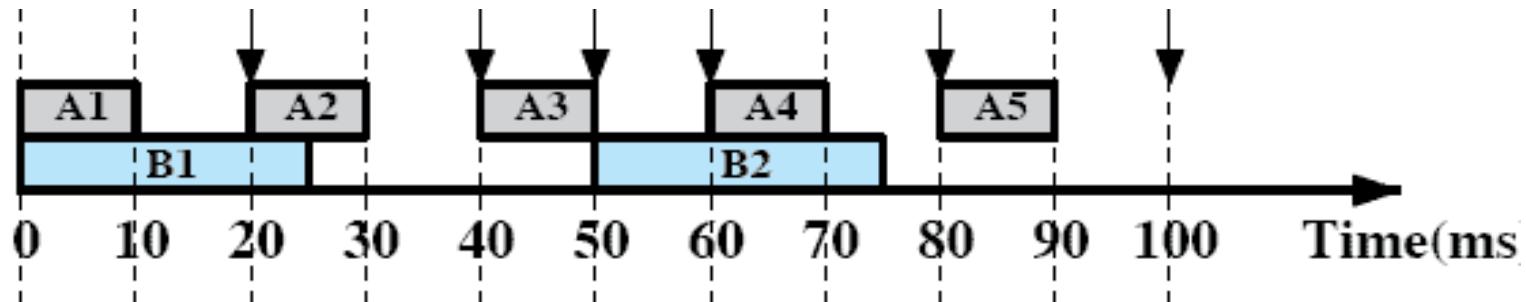
- A **hard real-time task** is one that must meet its deadline
 - Otherwise it will cause unacceptable damage or a fatal error to the system.

- A **soft real-time task** has an associated deadline that is desirable but not mandatory
 - it still makes sense to schedule and complete the task even if it has passed its deadline.

Periodic vs Aperiodic

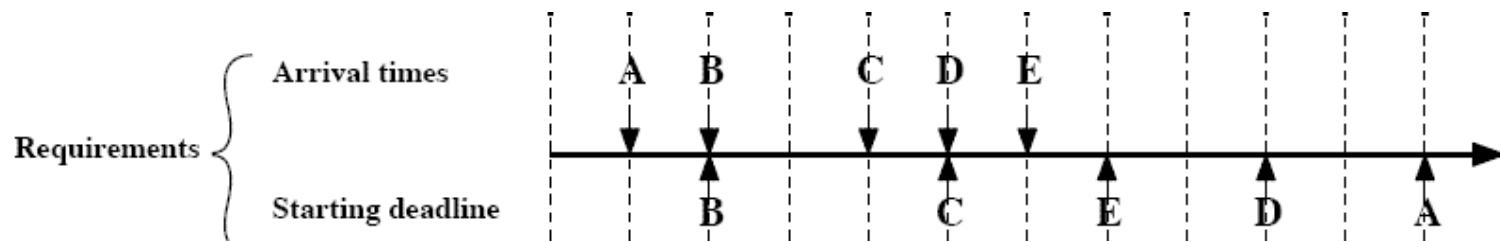
□ Periodic tasks

- The requirement may be stated as “once per period T ” or “exactly T units apart.”



□ Aperiodic tasks

- Have time constraints either for deadlines or start





Real-Time Systems

- Control of laboratory experiments
- Process control in industrial plants
- Air traffic control
- Telecommunications
- Military command and control systems
- Robotics



Features of Real-Time OS

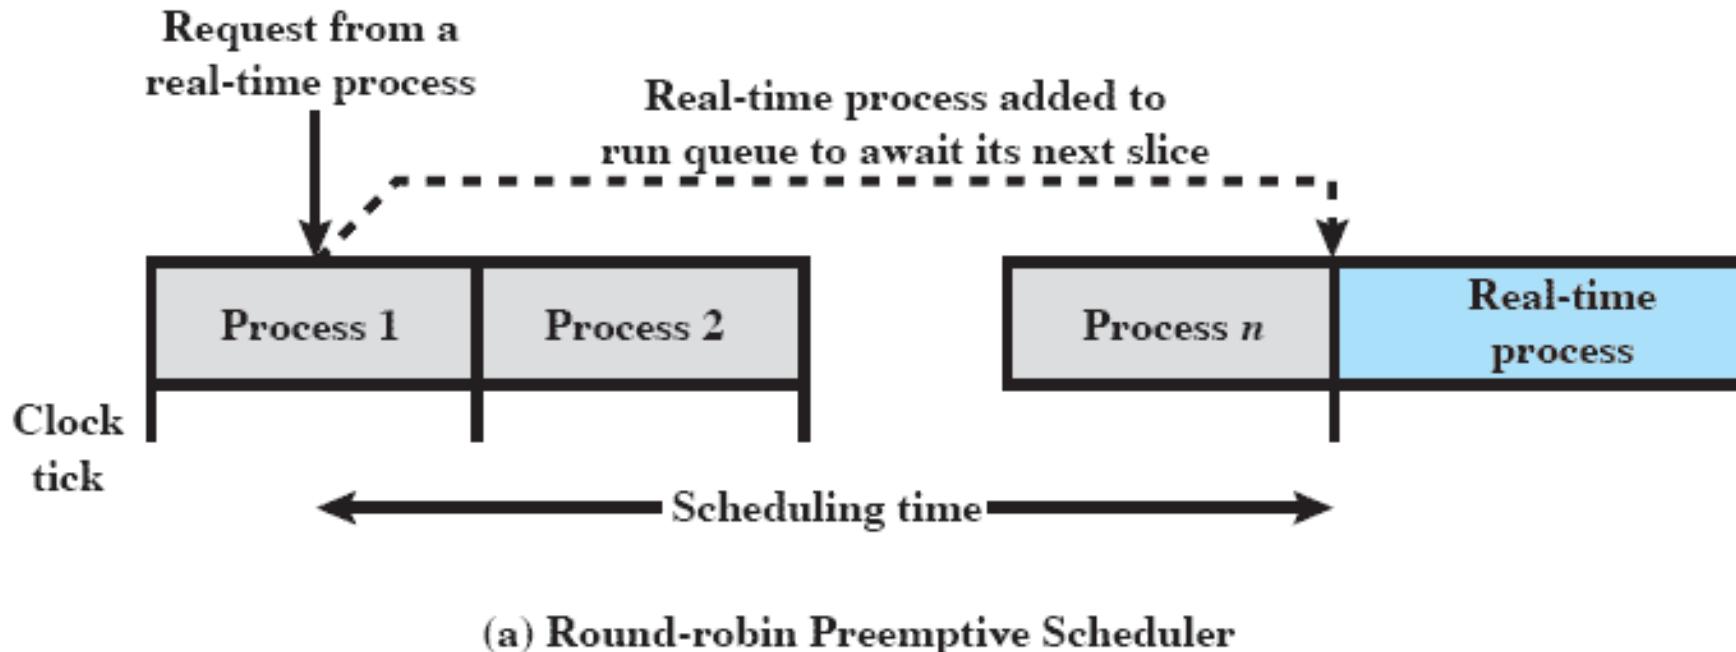
- Fast process or thread switch
- Small size
- Ability to respond to external interrupts quickly
- Multitasking with inter-process communication tools such as semaphores, signals, and events



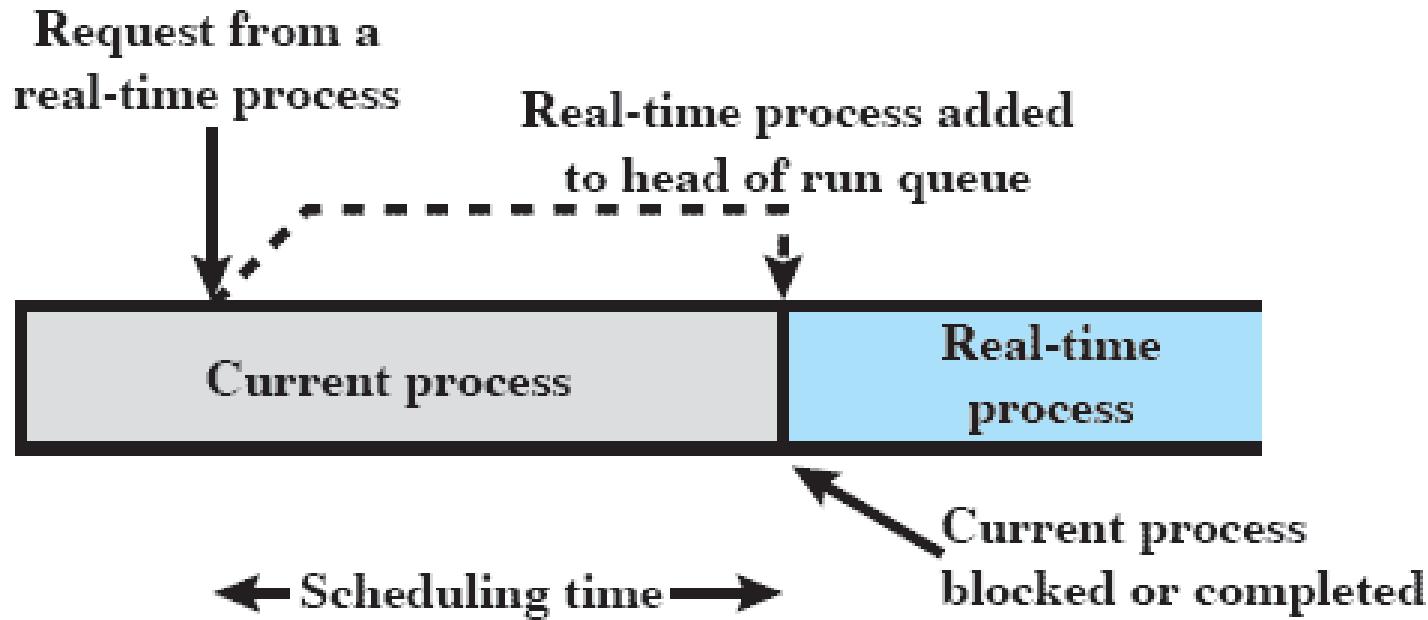
Features of Real-Time OS (cont...)

- Use of special sequential files that can accumulate data at a fast rate
- Preemptive scheduling based on priority
- Minimization of intervals during which interrupts are disabled
- Delay tasks for fixed amount of time
- Special alarms and timeouts

Round-robin Preemptive Scheduling: Unacceptable

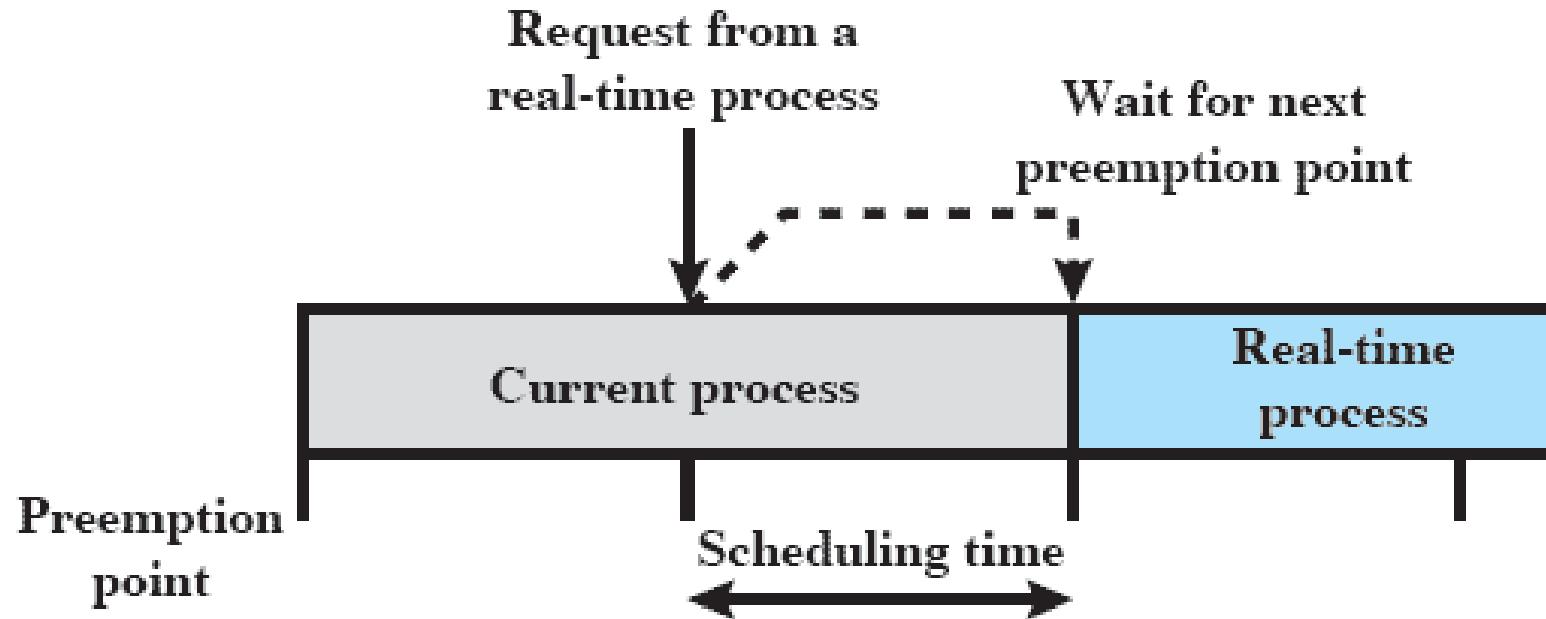


Priority-driven Nonpreemptive Scheduling: Unacceptable



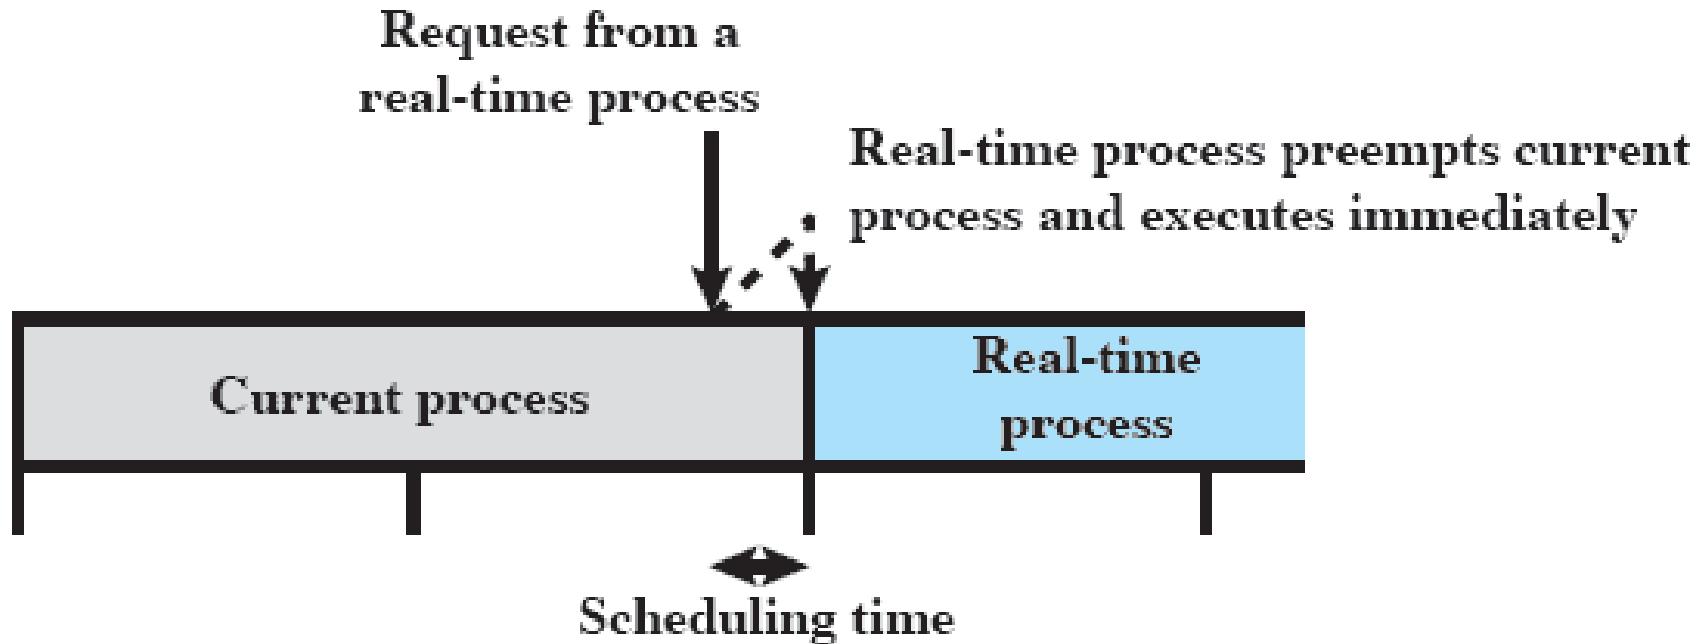
(b) Priority-Driven Nonpreemptive Scheduler

Combine Priorities with Clock-based Interrupts



(c) Priority-Driven Preemptive Scheduler on Preemption Points

Immediate Preemption



(d) Immediate Preemptive Scheduler



Classes of Real-Time Scheduling Algorithms

- Static table-driven
 - Performs a static analysis of feasible schedules of dispatching
 - Determines when a task must begin execution at run time
- Static priority-driven preemptive
 - A static analysis is performed, but no schedule is drawn up
 - Traditional priority-driven preemptive scheduler is used
- Dynamic planning-based
 - Feasibility determined at run time (dynamically) rather than offline
 - An arriving task is accepted for execution only if it is feasible to meet its time constraints
- Dynamic best effort
 - No feasibility analysis is performed
 - The system tries to meet all deadlines and aborts any started process whose deadline is missed



Policy: Deadline Scheduling

- Real-time applications are generally not concerned with sheer speed but with **completing (or starting) tasks**
 - at the most valuable times, neither too early nor too late
 - despite dynamic resource demands and conflicts, processing overloads, and hardware or software faults
- “**Priorities**” are a crude tool and may not capture the requirement of completion (or initiation) at the most valuable time



Deadline Scheduling

- Information used

- Ready time
- Starting deadline
- Completion deadline
- Processing time
- Resource requirements
- Priority
- Subtask scheduler

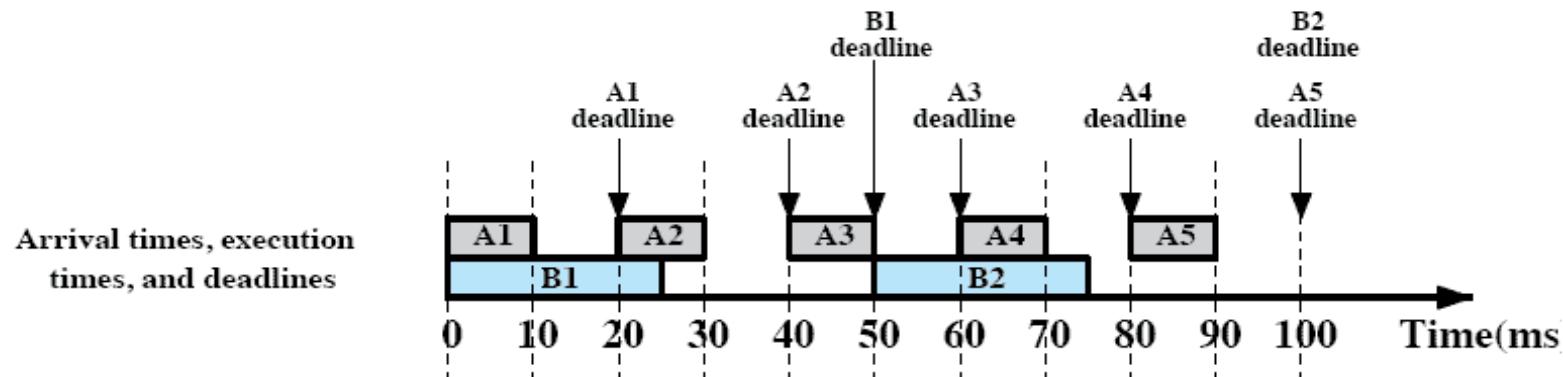
- For

- Periodic tasks
- Aperiodic tasks

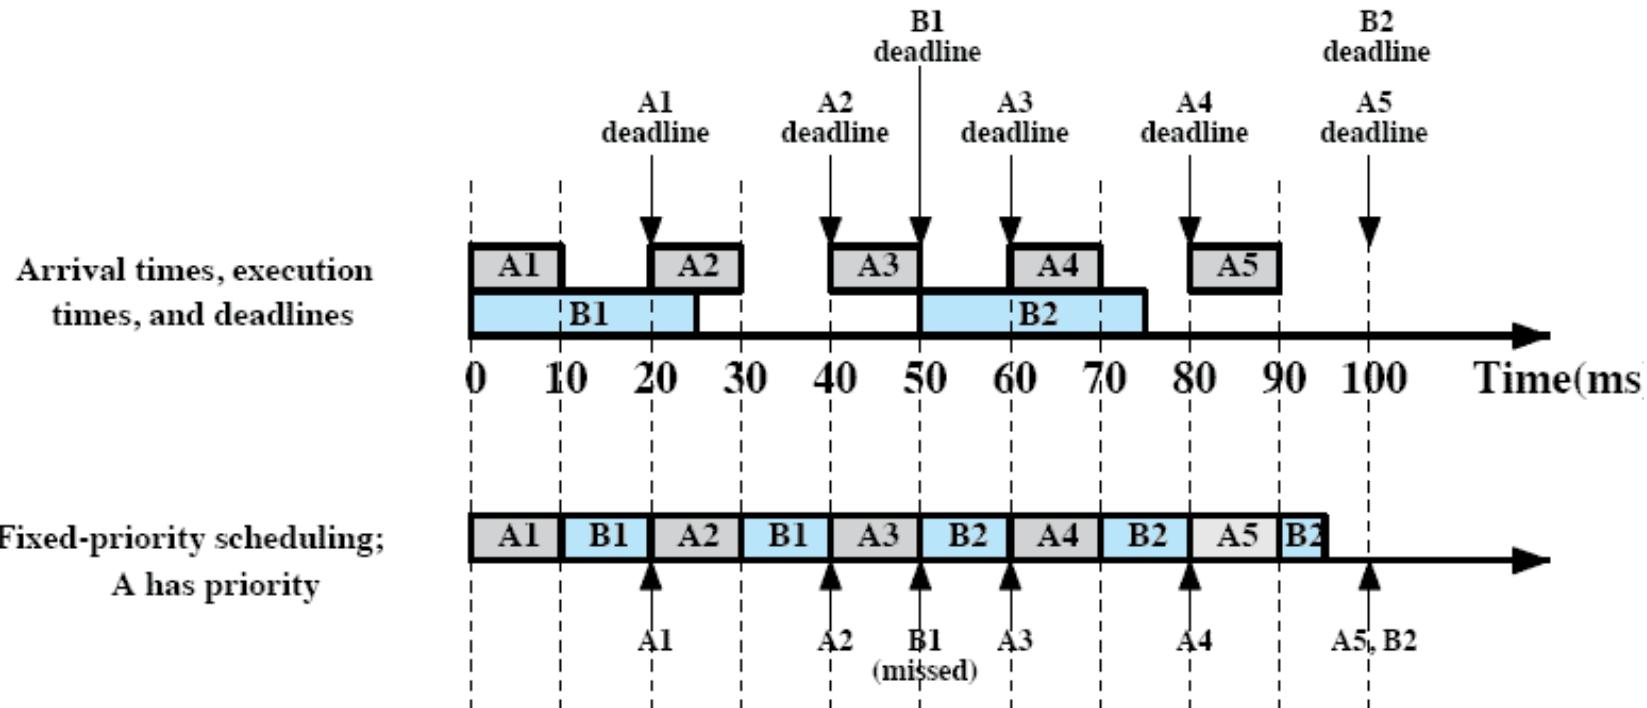
Periodic Scheduling: Two Periodic Tasks



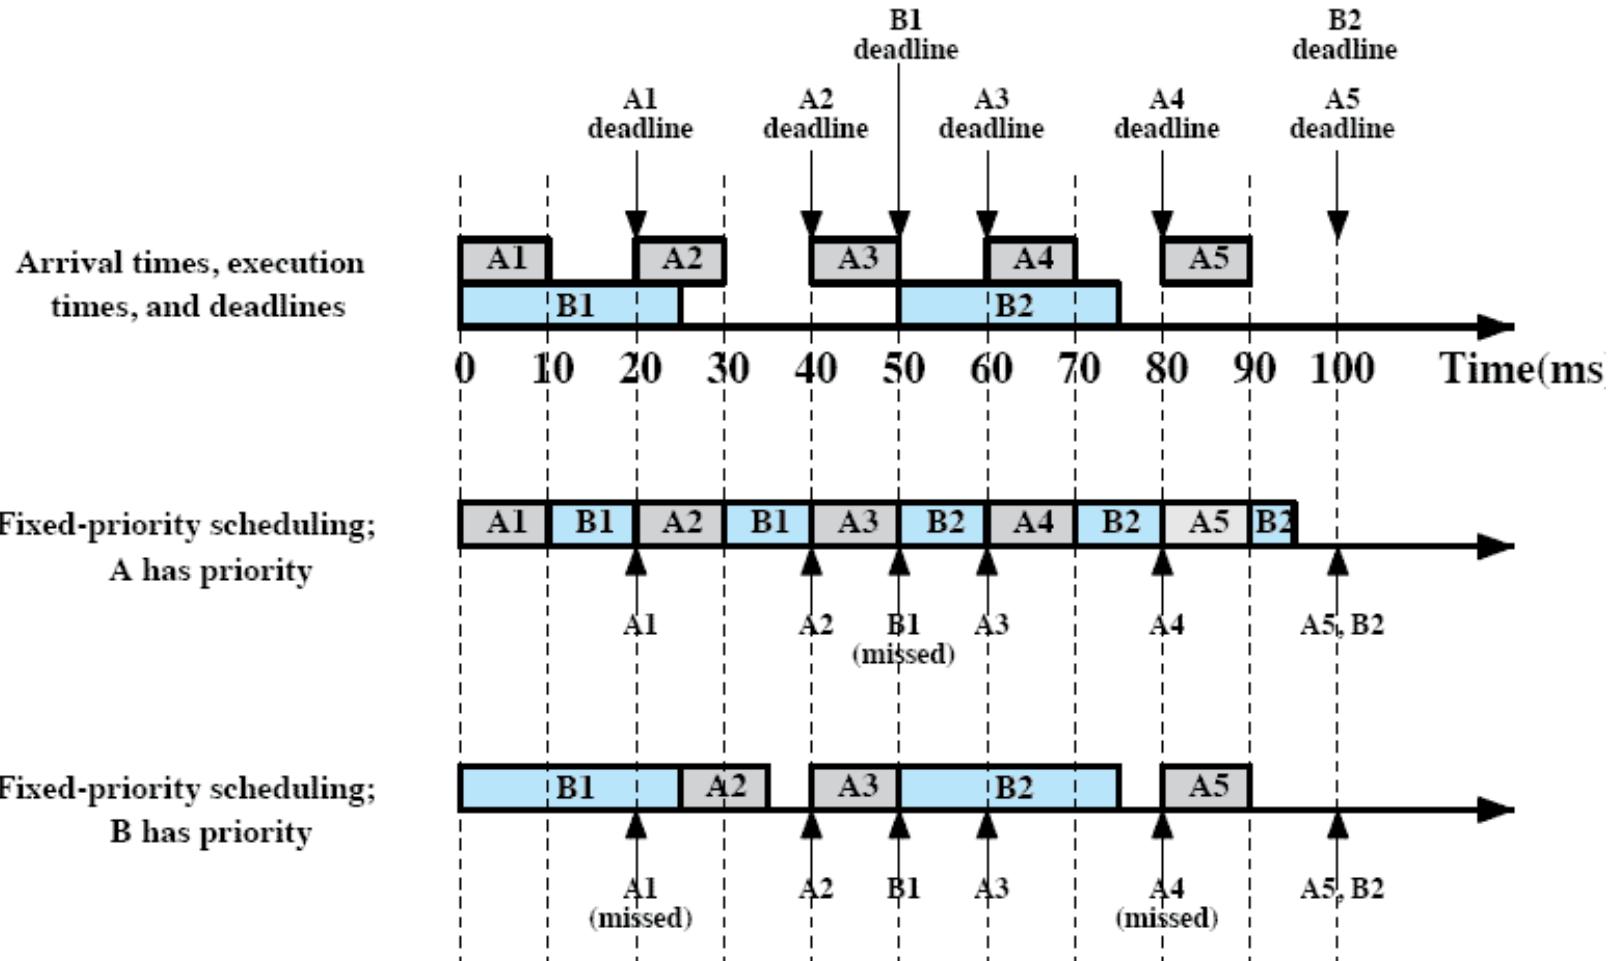
Process	Arrival Time	Execution Time	Ending Deadline
A(1)	0	10	20
A(2)	20	10	40
A(3)	40	10	60
A(4)	60	10	80
A(5)	80	10	100
⋮	⋮	⋮	⋮
⋮	⋮	⋮	⋮
⋮	⋮	⋮	⋮
B(1)	0	25	50
B(2)	50	25	100



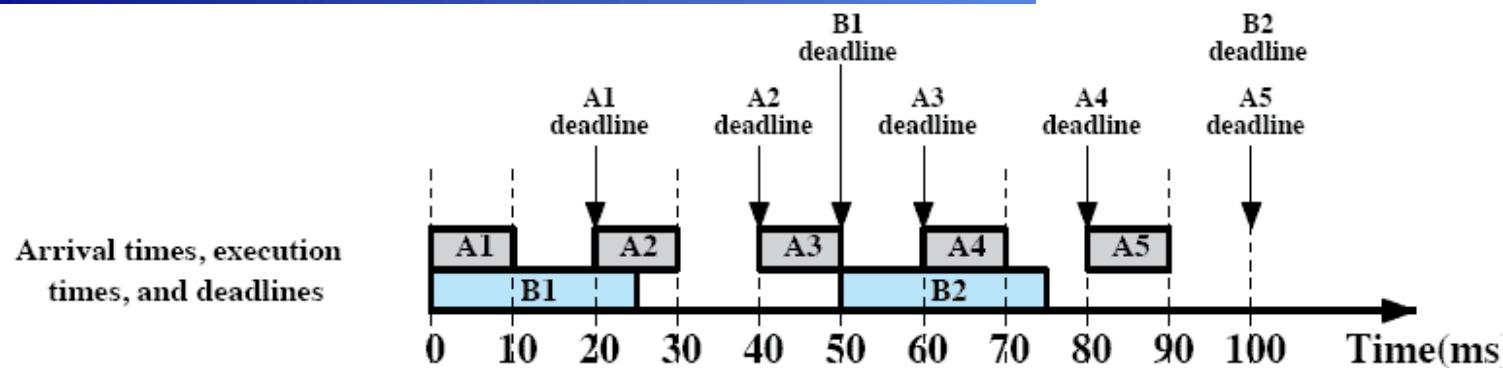
Periodic Scheduling



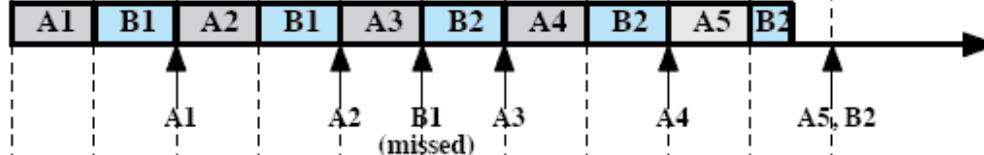
Periodic Scheduling



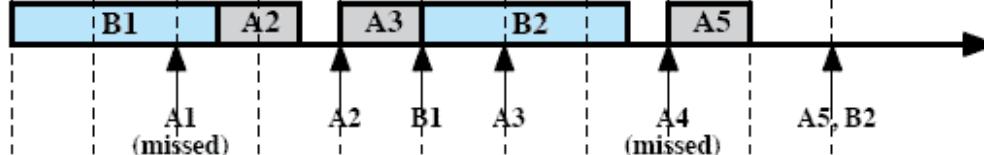
Periodic Scheduling



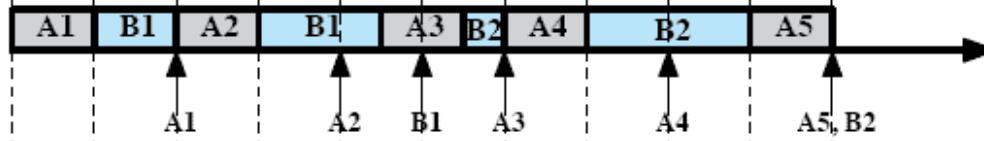
Fixed-priority scheduling;
A has priority



Fixed-priority scheduling;
B has priority



Earliest deadline scheduling
using completion deadlines



Periodic Scheduling: Priority-driven Preemptive Scheduling



Assumptions & Definitions

- Tasks are periodic
- No aperiodic or sporadic tasks
- Job (instance) deadline = end of period
- No resource constraints
- Tasks are preemptable

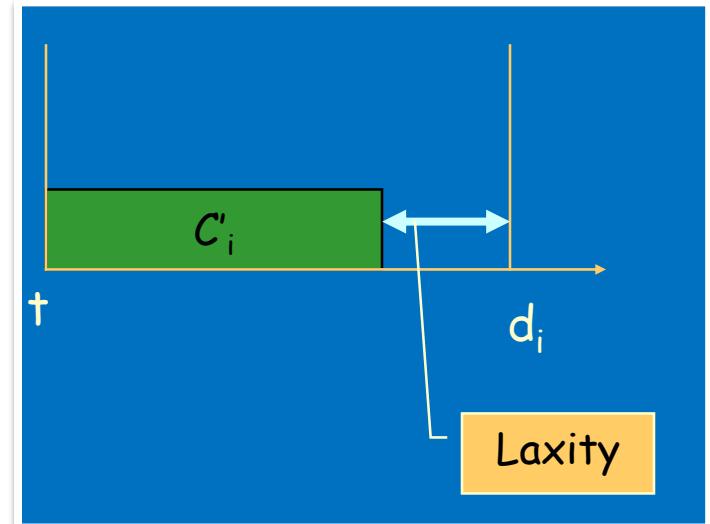
- Laxity of a Task

$$T_i = d_i - (t + c'_i)$$

where d_i : deadline;

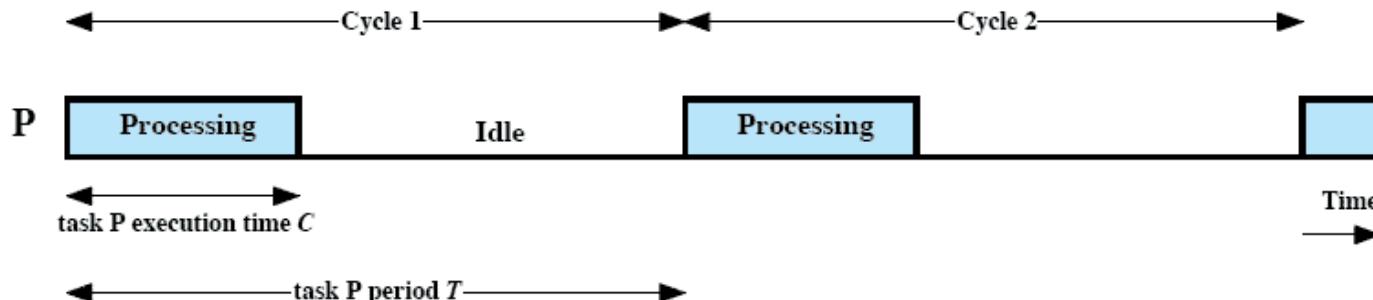
t : current time;

c'_i : remaining computation time.

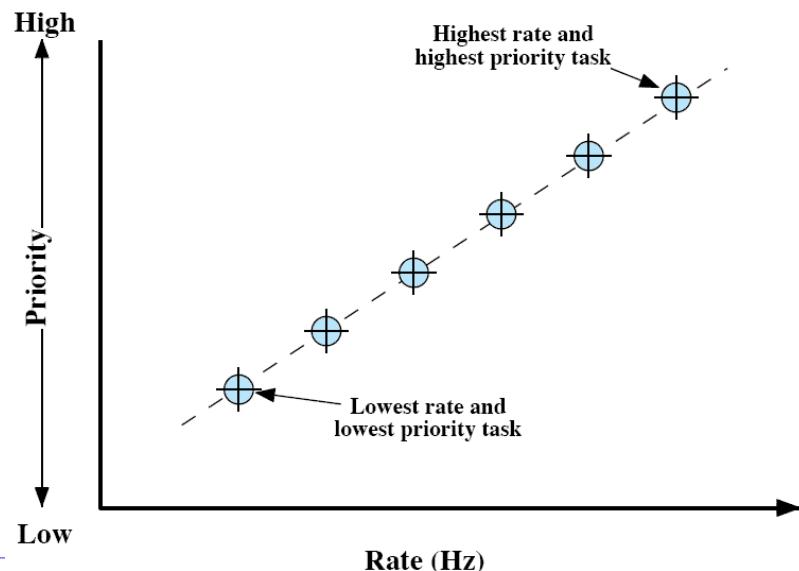


Rate Monotonic Scheduling (RMS)

- Assigns priorities to tasks on the basis of their periods



- The highest-priority task is the one with the shortest period



- The **rate monotonic algorithm (RMA)** is a procedure for assigning fixed priorities to tasks to maximize their **"schedulability."**
- A task set is considered **schedulable** if all tasks meet all deadlines all the time.



Rate Monotonic Scheduling (RMS)

□ Schedulability check (off-line)

- A set of n tasks is **schedulable** on a uniprocessor by the RMS algorithm if the processor utilization (utilization test):

$$\sum_{i=1}^n c_i/p_i \leq n(2^{1/n} - 1).$$

The term $n(2^{1/n} - 1)$ approaches $\ln 2$, (≈ 0.69 as $n \rightarrow \infty$).

- This condition is sufficient, but not necessary.



RMS (cont.)

- Schedule construction (online)
 - Task with the smallest period is assigned the highest priority.
 - At any time, the highest priority task is executed.
- RMS is an **optimal** preemptive scheduling algorithm with fixed priorities.
- Static/fixed priority algorithm assigns the same priority to all the jobs (instances) in each task.

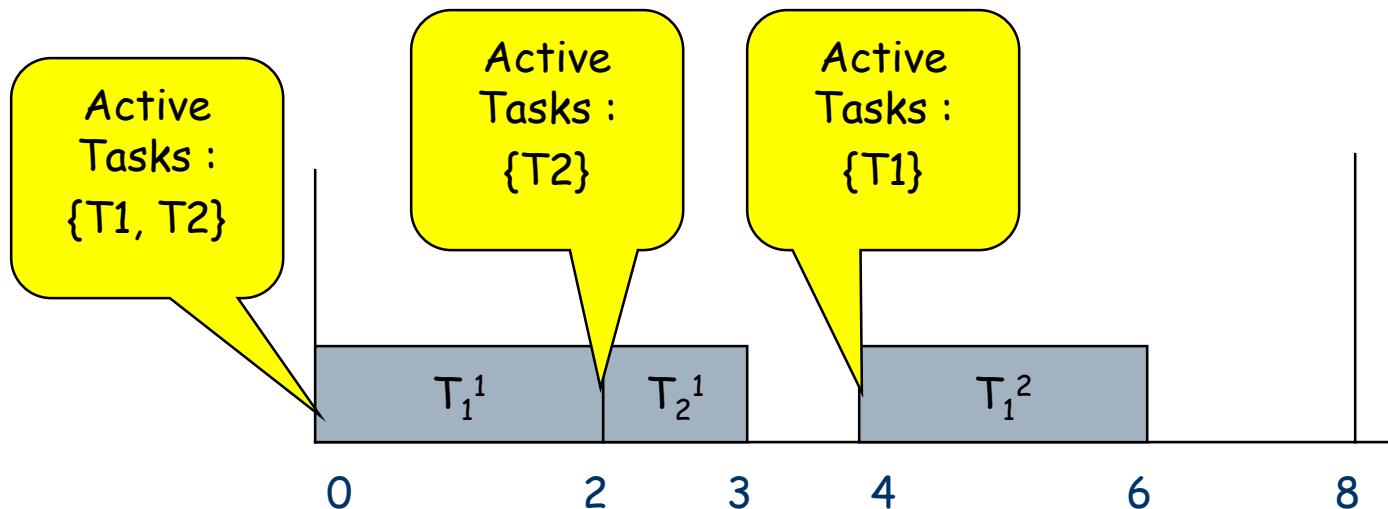
RMS Scheduler -- Example 1

Task set: $T_i = (c_i, p_i)$

$T1 = (2,4)$ and $T2 = (1,8)$

Schedulability check:

$$2/4 + 1/8 = 0.5 + 0.125 = 0.625 \leq 2(\sqrt{2} - 1) = 0.82$$



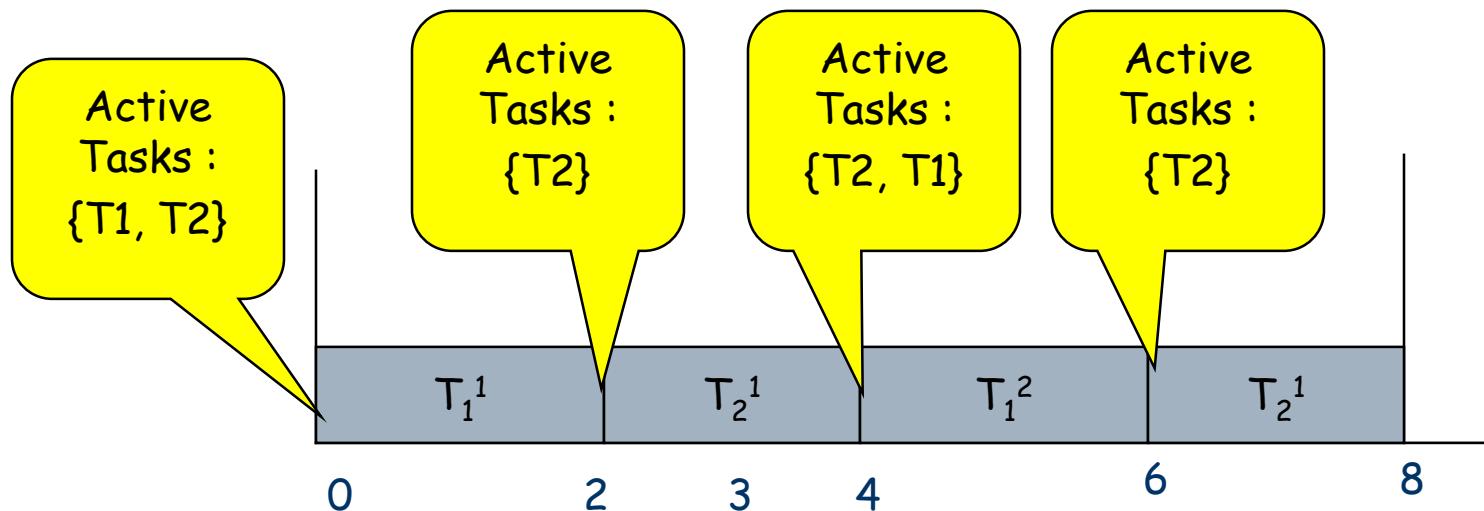
RMS scheduler -- Example-2

Task set: $T_i = (c_i, p_i)$

$T1 = (2,4)$ and $T2 = (4,8)$

Schedulability check:

$$2/4 + 4/8 = 0.5 + 0.5 = 1.0 > 2(\sqrt{2} - 1) = 0.82$$



Some task sets that FAIL the utilization-based schedulability test are also schedulable under RMS → We need exact analysis (necessary & sufficient)



Earliest Deadline First (EDF)

- **Schedulability check (off-line)**
 - A set of n tasks is schedulable on a uniprocessor by the EDF algorithm if the processor utilization

$$\sum_{i=1}^n c_i/p_i \leq 1$$

- This condition is both necessary and sufficient.
- Least Laxity First (LLF) algorithm has the same schedulability check.



EDF/LLF (cont.)

□ Schedule construction (online)

- EDF/LLF: Task with the smallest deadline/laxity is assigned the highest priority.
- At any time, the highest priority task is executed.

EDF/LLF is an optimal preemptive scheduling algorithm with dynamic priorities.

Dynamic priority algorithm assigns different priorities to the individual jobs (instances) in each task.

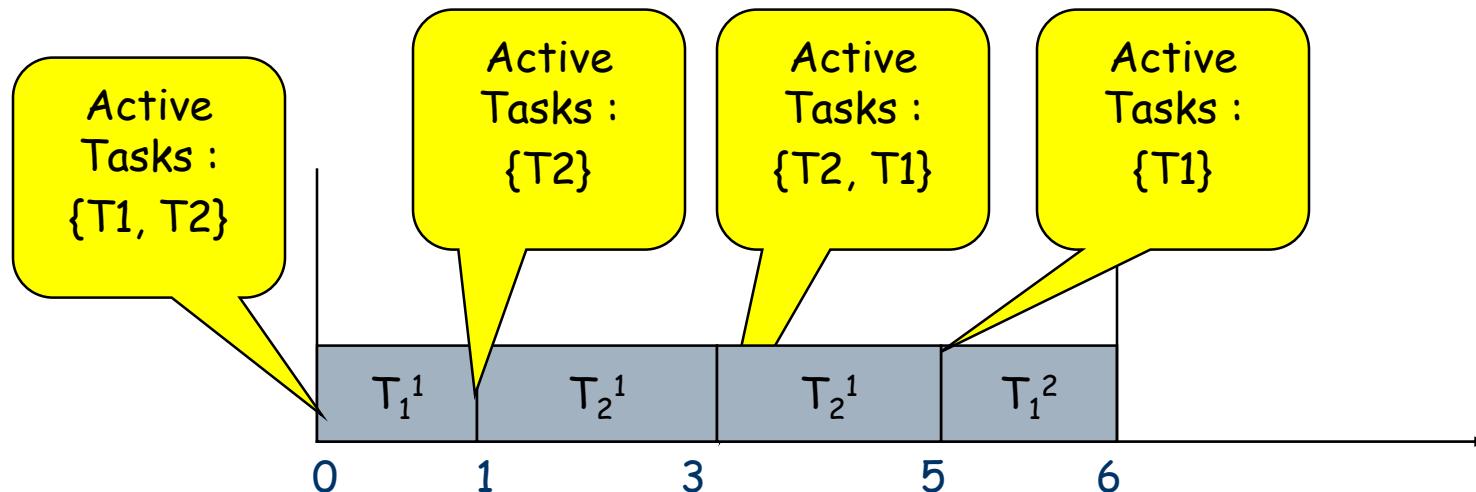
EDF scheduler -- Example

Task set: $T_i = (c_i, p_i, d_i)$

$T1 = (1,3,3)$ and $T2 = (4,6,6)$

Schedulability check:

$$1/3 + 4/6 = 0.33 + 0.67 = 1.0$$



Unlike RMS, Only those task sets which pass the schedulability test are schedulable under EDF



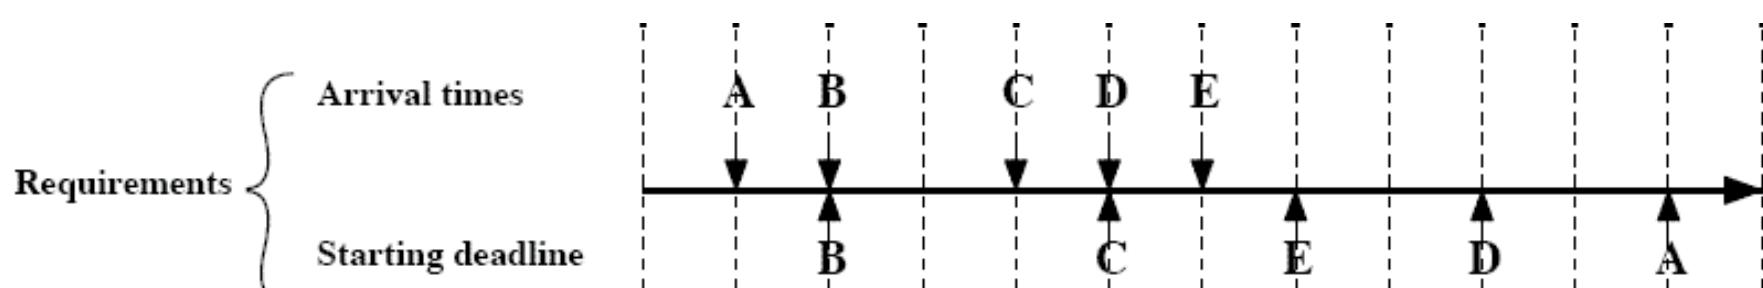
RMS vs. EDF/LLF

- RMS is an optimal preemptive scheduling algorithm with fixed priorities.
- EDF/LLF is an optimal preemptive scheduling algorithm with dynamic priorities.
- RMS schedulability properties can be analyzed; rich theory exists and it is widely used in practice.
- EDF/LLF offers higher schedulability than RMS, but it is more difficult to implement.

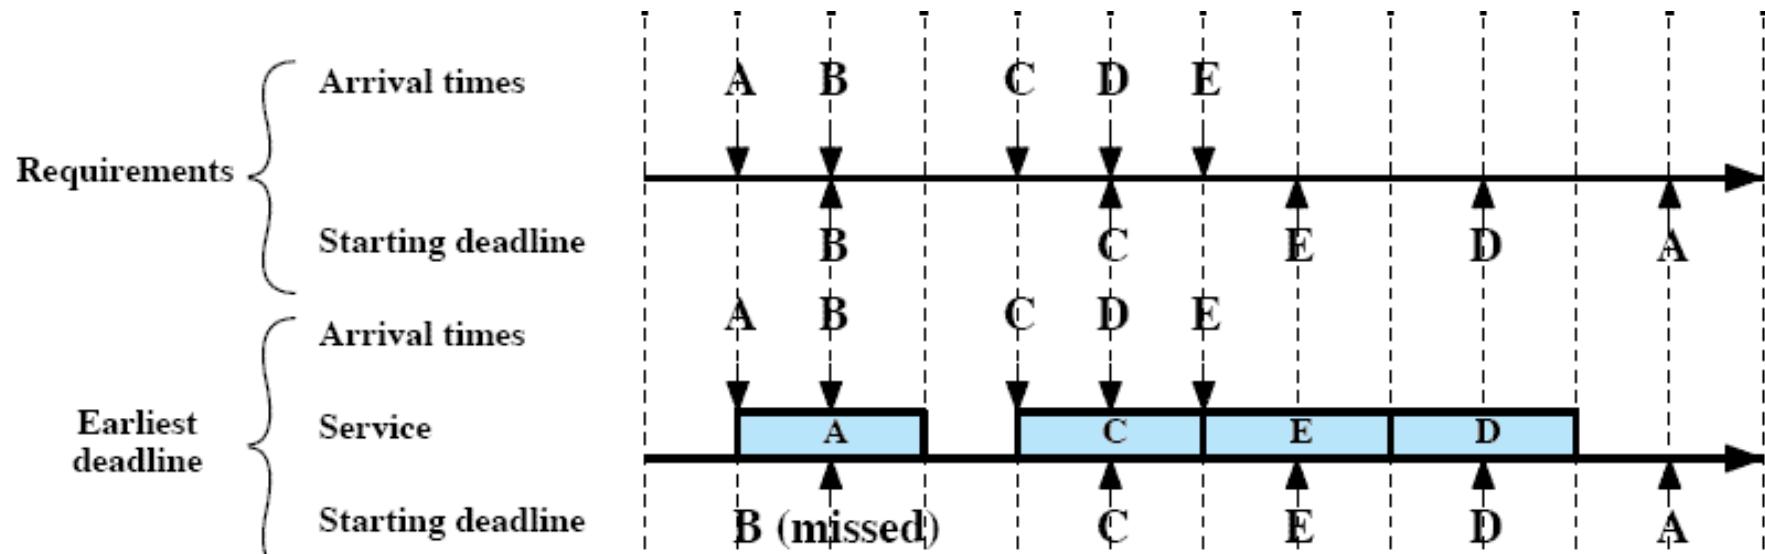
Aperiodic Scheduling: Five Aperiodic Tasks (Non-preemptive)



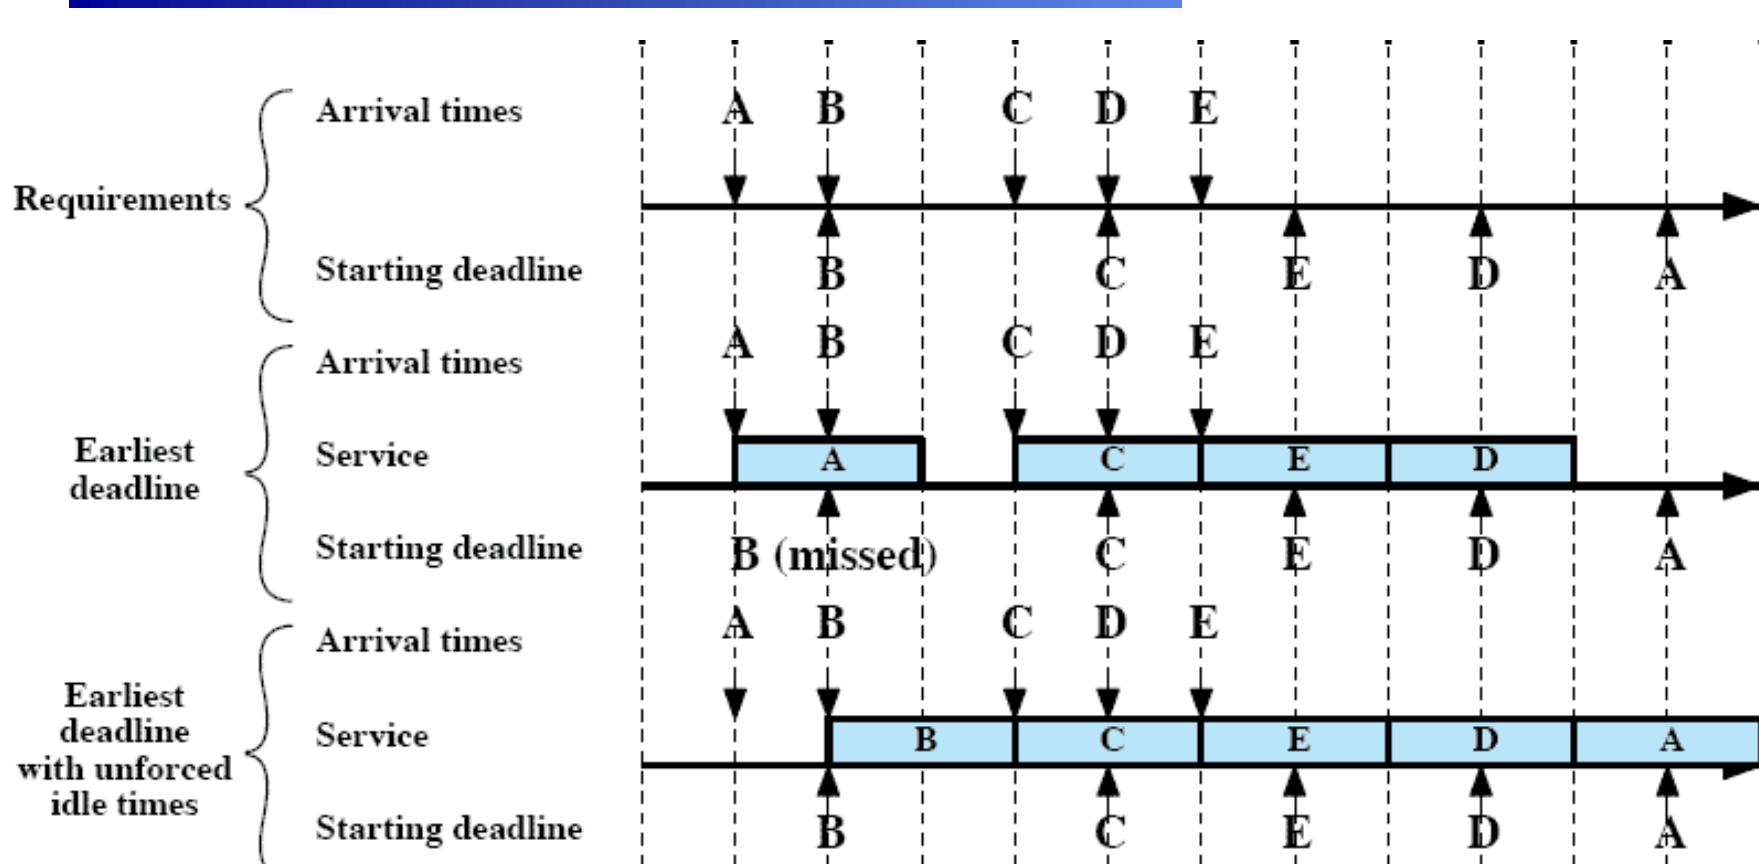
Process	Arrival Time	Execution Time	Starting Deadline
A	10	20	110
B	20	20	20
C	40	20	50
D	50	20	90
E	60	20	70



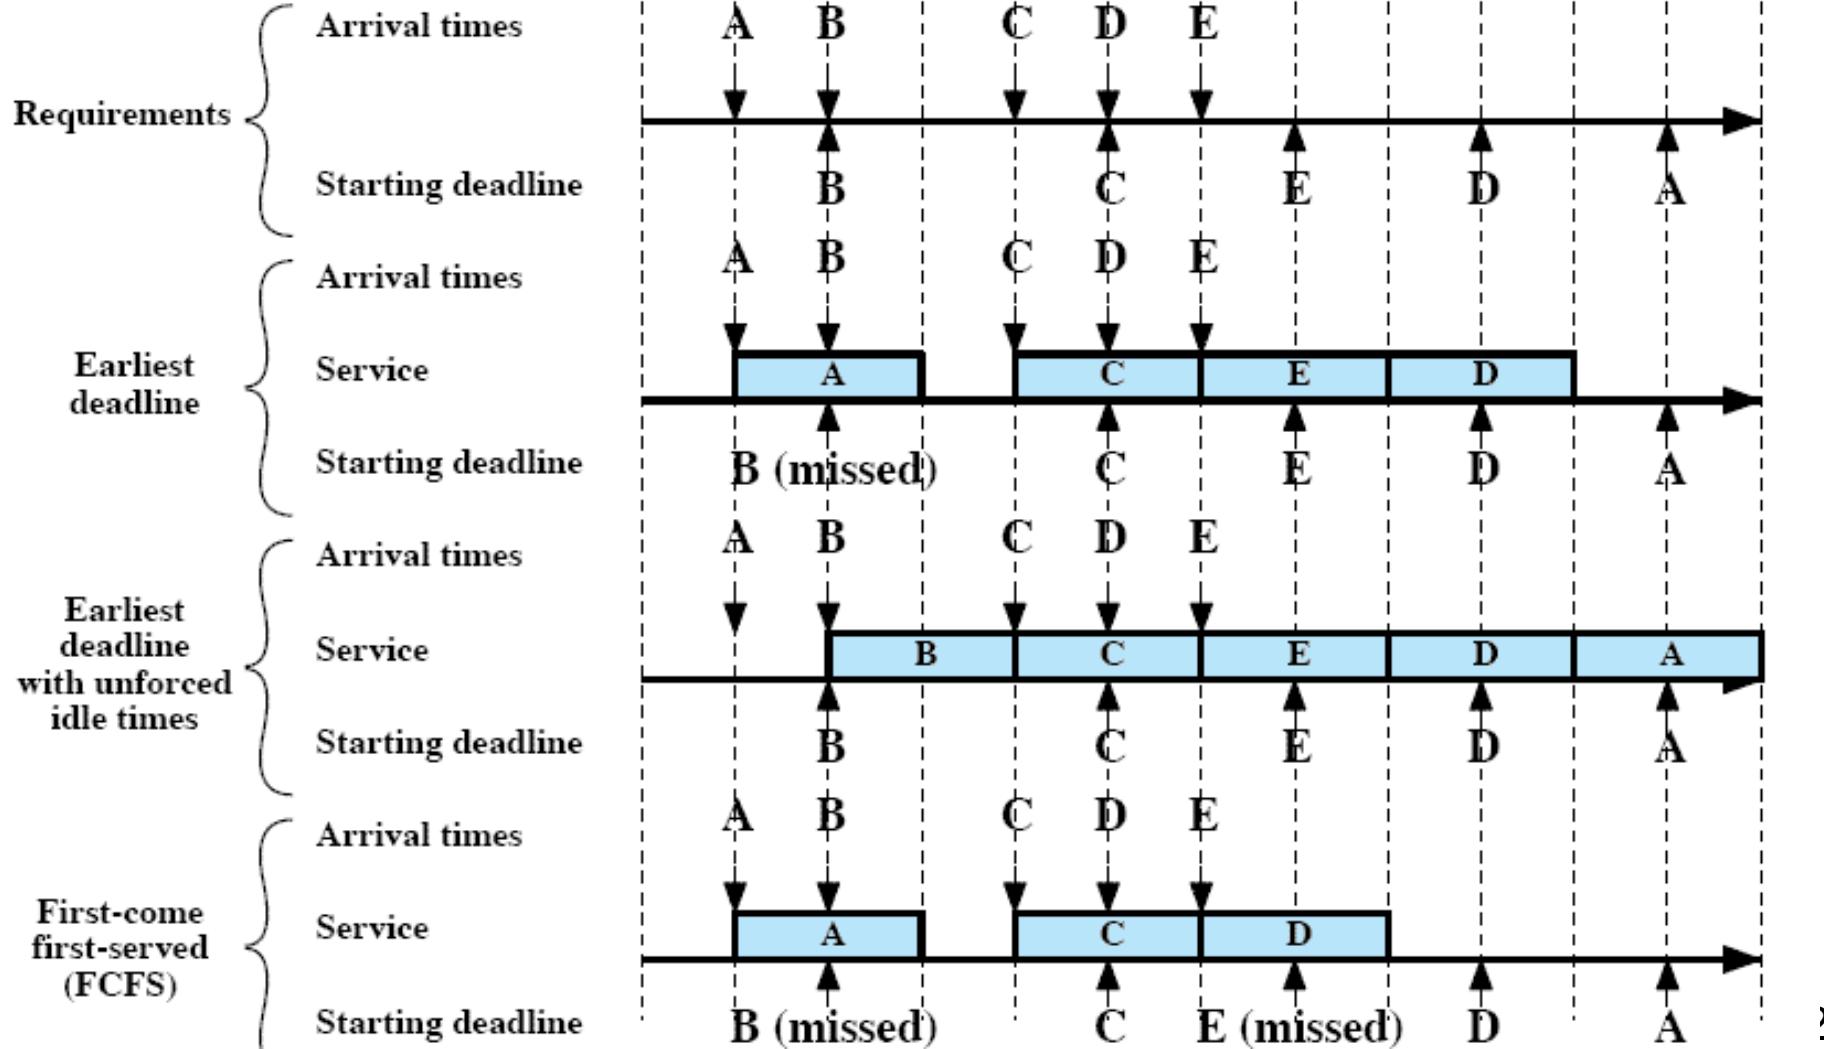
Aperiodic Scheduling



Aperiodic Scheduling



Aperiodic Scheduling





Homework: Scheduling

- Examine the scheduling policies in Linux and Windows and compare these policies (including policies for real-time scheduling).
 - Write a report, describe and compare the key scheduling policies in Linux and Windows.
 - How do they manage processes/threads?
 - What's the complexity (O) of their respective schedulers?
 - How do they handle priority inversion?
- Submission is not required



Summary

- Multiprocessor scheduling
- Real-time scheduling

- Next lecture
 - Synchronization (I) - Basics