



Operating Systems (A)

(Honor Track)

Lecture 15: Synchronization (II)

Yao Guo (郭耀)

Peking University

Fall 2021

Review



Synchronization Basics

Why do we need synchronizations
Critical sections



Mutual Exclusion

- We want to use **mutual exclusion** to synchronize access to shared resources
 - This allows us to have larger atomic blocks
- Code that uses mutual exclusion (and...) to synchronize its execution is called a **critical section**
 - Only one thread at a time can execute in the critical section
 - All other threads are forced to wait on entry
 - When a thread leaves a critical section, another can enter
- What requirements would you place on a critical section?



Critical Section Requirements

- Mutual exclusion (mutex)
 - If one thread is in the critical section, then no other is
- Progress
 - If some thread T is not in the critical section, then T cannot prevent some other thread S from entering the critical section
 - A thread in the critical section will eventually leave it
- Bounded waiting (no starvation)
 - If some thread T is waiting on the critical section, then T will eventually enter the critical section
- Performance
 - The overhead of entering and exiting the critical section is small with respect to the work being done within it

```
do {  
    entry section  
    critical section  
    exit section  
    remainder section  
} while (TRUE);
```

* Operating System Concept Essentials
(Abrahan Silberschatz et al.), P220



How to Build Critical Sections

- Mechanisms for building critical sections
 - Atomic read/write
 - Locks
 - Semaphores
 - Monitors
 - Messages



Peterson's Algorithm

```
int turn = 1;  
bool ready1 = false, ready2 = false;
```

```
while (true) {  
    ready1 = true;  
    turn = 2;  
    while (turn == 2 && ready2) ;  
    critical section  
    ready1 = false;  
    outside of critical section  
}
```

```
while (true) {  
    ready2 = true;  
    turn = 1;  
    while (turn == 1 && ready1) ;  
    critical section  
    ready2 = false;  
    outside of critical section  
}
```

- Does it work?
 - Mutual exclusion
 - Progress
 - Bounded waiting



Peterson's Algorithm (Figure 2-24)

```
#define FALSE 0
#define TRUE 1
#define N      2          /* number of processes */

int turn;                  /* whose turn is it? */
int interested[N];         /* all values initially 0 (FALSE) */

void enter_region(int process);    /* process is 0 or 1 */
{
    int other;               /* number of the other process */

    other = 1 - process;    /* the opposite of process */
    interested[process] = TRUE; /* show that you are interested */
    turn = process;          /* set flag */
    while (turn == process && interested[other] == TRUE) /* null statement */ ;
}

void leave_region(int process)    /* process: who is leaving */
{
    interested[process] = FALSE; /* indicate departure from critical region */
}
```



This Lecture

Common Synchronization Mechanisms

Locks

Semaphores

Monitors



Buzz Words

Lock

Spinlock

Semaphore

Starvation

Monitor

Condition variable



How to Build Critical Sections

- Mechanisms for building critical sections
 - Atomic read/write
 - Locks
 - Semaphores
 - Monitors
 - Messages



Locks

- A **lock** is an object in memory providing two operations
 - **acquire()**: before entering the critical section
 - **release()**: after leaving a critical section
 - Some systems may refer this pair as **Lock()** / **Unlock()**
- Threads **pair calls** to **acquire()** and **release()**
 - Between **acquire()/release()**, the thread **holds** the lock
 - **acquire()** does not return until any previous holder releases
 - What can happen if the calls are not paired?
- Locks can spin (a spinlock) or block (a mutex)



Using Locks

```
withdraw (account, amount) {  
    acquire(lock);  
    balance = get_balance(account);  
    balance = balance - amount;  
    put_balance(account, balance);  
    release(lock);  
    return balance;  
}
```

Critical
Section

```
acquire(lock);  
balance = get_balance(account);  
balance = balance - amount;
```

```
acquire(lock);
```

```
put_balance(account, balance);  
release(lock);
```

```
balance = get_balance(account);  
balance = balance - amount;  
put_balance(account, balance);  
release(lock);
```

- What happens when blue tries to acquire the lock?
- Why is the “return” outside the critical section? What if not?
- What happens when a third thread calls acquire?



Implementing Locks (1/3)

- How do we implement locks? Here is one attempt:

```
struct lock {  
    int held = 0;  
}  
  
void acquire (lock) {  
    while (lock->held);  
    lock->held = 1;  
}  
  
void release (lock) {  
    lock->held = 0;  
}
```

busy-wait (spin-wait)
for lock to be released

- This is called a **spinlock** because a thread spins waiting for the lock to be released
- Does this work?



Implementing Locks (2/3)

- No. Two independent threads may both notice that a lock has been released and thereby acquire it

```
struct lock {  
    int held = 0;  
}  
  
void acquire (lock) {  
    while (lock->held);  
    lock->held = 1;  
}  
  
void release (lock) {  
    lock->held = 0;  
}
```

A context switch can occur here, causing a race condition



Implementing Locks (3/3)

- The problem is that the implementation of locks has critical sections, too
 - How do we stop the recursion?
- The implementation of acquire/release must be **atomic**
- How do we make them atomic?
- Need help from hardware
 - Atomic instructions (e.g., **test-and-set**)
 - Disable/enable interrupts (prevents context switches)



Atomic Instructions: Test-And-Set

- The semantics of test-and-set are:
 - Record the old value
 - Set the value to TRUE
 - Return the old value
- Hardware executes it **atomically!**

```
bool test_and_set (bool *flag) {  
    bool old = *flag;  
    *flag = True;  
    return old;  
}
```

- When executing test-and-set on “flag”
 - What is **value of flag** afterwards if it was initially False? True?
 - What is the **return result** if flag was initially False? True?



Using Test-And-Set

- Here is our lock implementation with **test-and-set**:

```
struct lock {  
    int held = 0;  
}  
void acquire (lock) {  
    while (test-and-set(&lock->held));  
}  
void release (lock) {  
    lock->held = 0;  
}
```

- When will the **while** return? What is the value of **held**?
- Does it work?
- Does it work on multiprocessors?



Other Similar Hardware Instruction

□ Swap/XCHG

```
void Swap (char* x,* y);
{ // All done atomically
    char temp = *x;
    *x = *y;
    *y = temp
}
```

□ How to implement Test-And-Set using Swap/XCHG?

```
bool test_and_set (bool *flag)
{
    bool X=True;
    Swap(&X, &flag);
    return X;
}
```



Problems with Spinlocks

- The problem with spinlocks is that they are wasteful
 - If a thread is spinning on a lock, it occupies the CPU, and the other thread holding the lock is not executing to get out of the critical section and release the lock
- Solution:
 - If a thread cannot get the lock, call `thread_yield` to give up the CPU!
- Solution 2: sleep and wakeup
 - When blocked, go to sleep
 - Wakeup when it is OK to retry entering the critical section



Disabling Interrupts

- Another implementation of `acquire/release` is to `disable interrupts`:

```
struct lock {  
}  
  
void acquire (lock) {  
    disable interrupts;  
}  
  
void release (lock) {  
    enable interrupts;  
}
```

- Note that there is no state associated with the lock
 - There is nothing saying you are locked or not; no `lock->held`
- Can two threads disable interrupts simultaneously in one processor core?
- Does it work on multi-processors?



On Disabling Interrupts

- Disabling interrupts blocks notification of external events that could trigger a context switch (e.g., timer)

- Only use disabling interrupts to implement higher-level synchronization primitives
 - Don't want interrupts disabled between acquire and release



Issue with the current solutions

- Goal: Use **mutual exclusion** to protect **critical sections** of code that access **shared resources**
- Method: Use locks (spinlocks or disable interrupts)
- Problem: Critical sections can be long

Spinlocks:

- Threads waiting to acquire lock spin in test-and-set loop
- Wastes CPU cycles
- Longer the CS, the longer the spin
- Longer the CS, greater the chance for lock holder to be interrupted

```
{  
    acquire(lock)  
    ...  
    Critical section  
    ...  
    release(lock)  
}
```

Disabling Interrupts:

- Should not disable interrupts for long periods of time
- Can miss or delay important events (e.g., timer, I/O)

Do we have any solution?

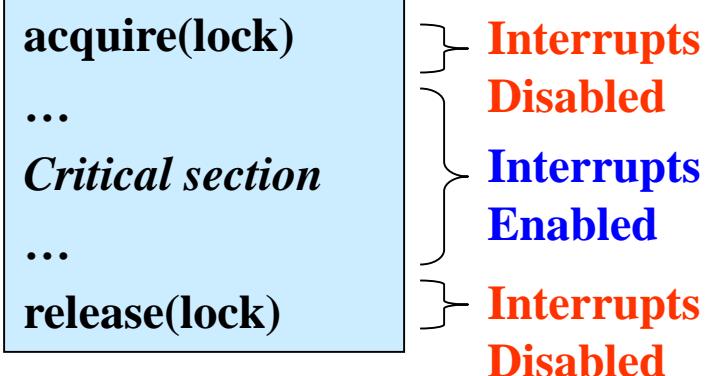


Solution

- Implementing locks with interrupts enabled in critical sections

```
struct lock {  
    int held = 0;  
    queue Q;  
}  
  
void acquire (lock) {  
    Disable interrupts;  
    while (lock->held) {  
        put current thread on lock Q;  
        Enable interrupts;  
        block current thread;  
        Disable interrupts;  
    }  
    lock->held = 1;  
    Enable interrupts;  
}
```

```
void release (lock) {  
    Disable interrupts;  
    if (Q) remove waiting thread;  
    unblock waiting thread;  
    lock->held = 0;  
    Enable interrupts;  
}
```





Another Issue of Disabling Interrupts

- Disabling interrupts is insufficient on a multiprocessor
- So we are back to atomic instructions



This Lecture

Common Synchronization Mechanisms

Locks

Semaphores

Monitors



Higher-Level Synchronization

- We looked at using locks to provide mutual exclusion
- Those locks work, but they have some drawbacks when critical sections are long
 - Spinlocks – inefficient
 - Disabling interrupts – can miss or delay important events, and do not work on multicores.
- Instead, we want synchronization mechanisms that
 - Block waiters
 - Leave interrupts enabled inside the critical section
- Look at some common high-level mechanism
 - **Semaphores**: binary (mutex) and counting



Semaphores

- Semaphores are an **abstract data type** that provides mutual exclusion to critical sections
- Semaphores are **integers** that support two operations:
 - **P(semaphore)**: decrement, block until semaphore is open
 - Also Wait(), or Down()
 - **V(semaphore)**: increment, allow another thread to enter
 - Also Signal(), or Up()
 - That's it! No other operations – not even just reading its value – exist
- Semaphore **safety property**: the semaphore value is always greater than or equal to 0
- Semaphores can also be used as atomic counters

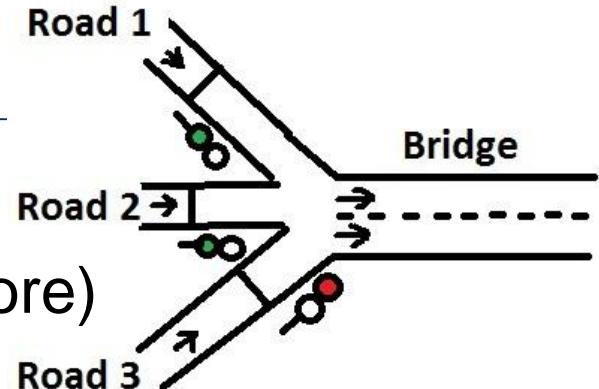


Blocking in Semaphores

- Associated with each semaphore is a queue of waiting processes
- When `P()` is called by a thread:
 - If semaphore is open, thread continues
 - If semaphore is closed, thread blocks on queue
- Then `V()` opens the semaphore:
 - If a thread is waiting on the queue, the thread is unblocked
 - If no threads are waiting on the queue, the signal is remembered for the next thread
 - In other words, `V()` has “history”
 - This “history” is a counter

Semaphore Types

- Semaphores come in two types
- **Mutex** semaphore (or **binary** semaphore)
 - Represents single access to a resource
 - Guarantees mutual exclusion to a critical section
- **Counting** semaphore (or **general** semaphore)
 - Represents a resource with many units available, or a resource that allows certain kinds of unsynchronized concurrent access (e.g., reading)
 - **Multiple threads can pass the semaphore**
 - Number of threads determined by the semaphore “count”
 - mutex has count = 1, counting has count = N
- You can use one type to implement the other
- Does a counting semaphore implement a critical section?



B. Komazec, 2012



Semaphore Semantics

- NOT the semaphore's implementation
 - They should be **atomic**
 - Any better idea than busy waiting?

```
void P (int *semaphore) {  
    while (*semaphore<=0);  
    *semaphore--;  
}
```

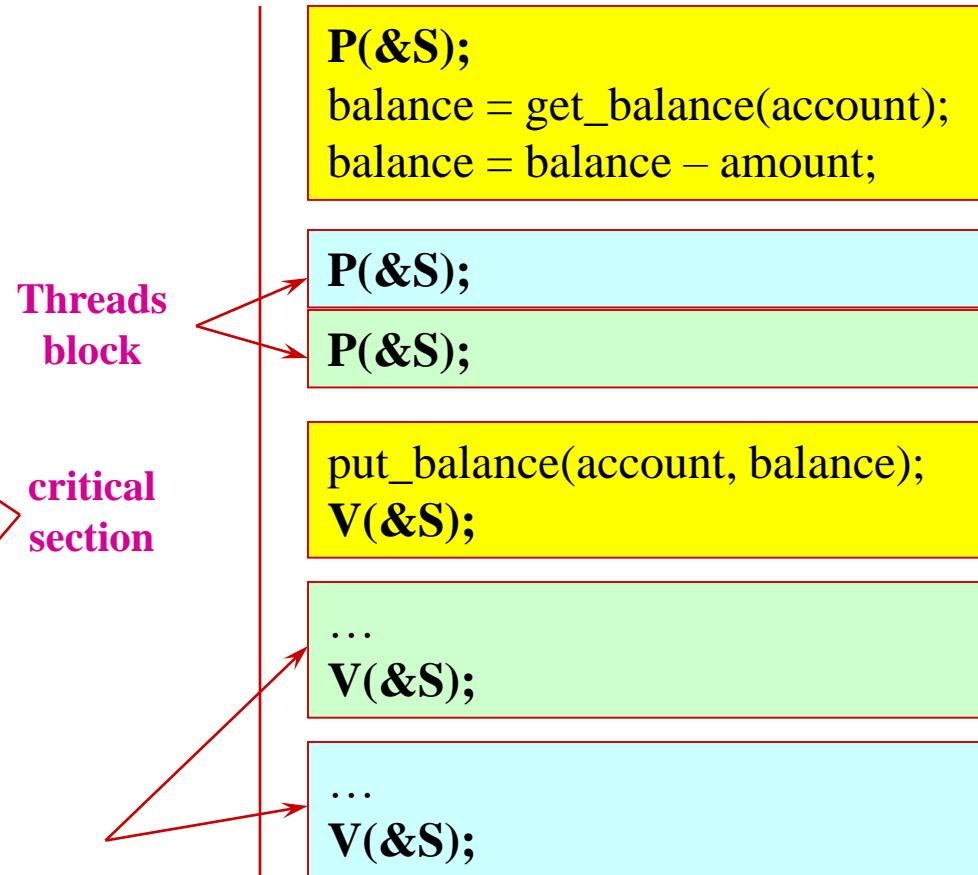
```
void V (int *semaphore) {  
    *semaphore++;  
}
```

Using Semaphores

- ❑ Mutex is similar to our locks, but semantics are different

```
struct Semaphore {
    int value;
    Queue q;
} S;

withdraw (account, amount) {
    ...
    P(&S);
    balance = get_balance(account);
    balance = balance - amount;
    put_balance(account, balance);
    V(&S);
    return balance;
}
```



**It is undefined which thread
runs after the V**



Possible Deadlocks with Semaphores

Example:

P0

P1

share two mutex semaphores S and Q

S:= 1; Q:=1;

P(S);

P(Q);

.....

V(S);

V(Q);

P(Q);

P(S);

.....

V(Q);

V(S);

Deadlock?



Be Careful When Using Semaphores

// Violation of Mutual Exclusion

V(mutex);

critical section

P(mutex);

// Deadlock Situation

P(mutex);

critical section

P(mutex);

// Violation of Mutual Exclusion

critical section

V(mutex);



A Classic Synchronization Problem

- We've looked at a simple example for using synchronization
 - Mutual exclusion while accessing a bank account

- Now we're going to use semaphores to look at a more interesting example
 - Readers/Writers



Readers/Writers Problem

- http://en.wikipedia.org/wiki/Readers-writers_problem
- Readers/Writers Problem:
 - An object is shared among several threads
 - Some threads only read the object, others only write it
 - We can allow **multiple readers** but only **one writer**
- How can we use semaphores to control access to the object to implement this protocol?



First Attempt

```
Semaphore w_or_r=1;

Reader{
    P(w_or_r); // lock out writers
    read;
    V(w_or_r); // up for grabs
}

writer {
    P(w_or_r); // lock out readers
    write;
    V(w_or_r); // up for grabs
}
```

Does it work?

Why?



Second Attempt

```
Semaphore w_or_r=1;  
int readcount; //record #readers  
  
Reader{  
    readcount++;  
    if (readcount == 1){  
        P(w_or_r); // lock out writers  
    }  
    read;  
    readcount--;  
    if (readcount == 0){  
        V(w_or_r); // up for grabs  
    }  
}  
  
writer {  
    P(w_or_r); // lock out readers  
    write;  
    V(w_or_r); // up for grabs  
}
```

Does it work?

Why?

Is **readcount** a shared data?
Who protects it?



Readers/Writers: Real Solution

□ Use three variables

- int **readcount** – number of threads reading object
- Semaphore **mutex** – control access to readcount
- Semaphore **w_or_r** – exclusive writing or reading

```
// number of readers
int readcount = 0;
// mutual exclusion to readcount
Semaphore mutex = 1;
// exclusive writer or reader
Semaphore w_or_r = 1;
```

```
writer {
    P(w_or_r); // lock out readers
    Write;
    V(w_or_r); // up for grabs
}
```

```
reader {
    P(mutex); // lock readcount
    readcount++; // one more reader
    if (readcount == 1)
        P(w_or_r); // synch w/ writers
    V(mutex); // unlock readcount
    Read;
    P(mutex); // lock readcount
    readcount--; // one less reader
    if (readcount == 0)
        V(w_or_r); // up for grabs
    V(mutex); // unlock readcount
}
```



Readers/Writers Notes

- `w_or_r` provides mutex between readers and writers, and also between multiple writers
- Why do readers use mutex (binary semaphore)?
- What if `V()` is above “`if (readcount == 1)`”?
- Why do we need “`if(readcount==1)`”?
- Why do we need “`if(readcount==0)`”?

```
reader {  
    P(mutex);      // lock readcount  
    readcount++; // one more reader  
    if (readcount == 1)  
        P(w_or_r); // synch w/ writers  
    V(mutex); // unlock readcount  
    Read;  
    P(mutex);      // lock readcount  
    readcount--; // one less reader  
    if (readcount == 0)  
        V(w_or_r); // up for grabs  
    V(mutex); // unlock readcount}  
}
```



But It Still Has a Problem

- What if a writer is waiting, but readers keep coming?

```
// number of readers
int readcount = 0;
// mutual exclusion to readcount
Semaphore mutex = 1;
// exclusive writer or reader
Semaphore w_or_r = 1;

writer {
    P(w_or_r); // lock out readers
    Write;
    V(w_or_r); // up for grabs
}
```

```
reader {
    P(mutex); // lock readcount
    readcount++; // one more reader
    if (readcount == 1)
        P(w_or_r); // synch w/ writers
    V(mutex); // unlock readcount
    Read;
    P(mutex); // lock readcount
    readcount--; // one less reader
    if (readcount == 0)
        V(w_or_r); // up for grabs
    V(mutex); // unlock readcount}
```

Problem: Starvation

- What if a writer is waiting, but readers keep coming, the writer is **starved!**



- How to write a solution without starvation?



Thoughts on Semaphores

- Semaphores can be used to solve the traditional synchronization problems
- However, they have some drawbacks
 - They are essentially **shared global variables**
 - Can potentially be accessed anywhere in program
 - No connection between the semaphore and the data being controlled by the semaphore
 - Used both for critical sections (mutual exclusion) and coordination (scheduling)
 - Note that I had to use comments in the code to distinguish
 - No control or guarantee of proper usage
- Sometimes hard to use and prone to bugs



This Lecture

Common Synchronization Mechanisms

Locks

Semaphores

Monitors



Monitors

- A monitor is a programming language construct that controls access to shared data
 - Synchronization code added by compiler, enforced at runtime
 - Why is this an advantage?
- A monitor is a module that encapsulates
 - Shared data structures
 - Procedures that operate on the shared data structures
 - Synchronization between concurrent threads that invoke the procedures
- A monitor protects its data from unsynchronized access
- It guarantees that threads accessing its data through its procedures interact only in legitimate ways

```
Monitor account {  
    double balance;  
  
    double withdraw(amount) {  
        balance = balance - amount;  
        return balance;  
    }  
}
```



Monitor Semantics

- A monitor guarantees “mutual exclusion?”
 - Only one thread can execute any monitor procedure at any time (the thread is “in the monitor”)
 - If a second thread invokes a monitor procedure when a first thread is already executing one, it blocks
 - So the monitor has to have a wait queue...
 - **If a thread within a monitor blocks, another one can enter**
- What’s the difference with critical sections we’ve learnt?
- What are the implications in terms of parallelism in monitor?

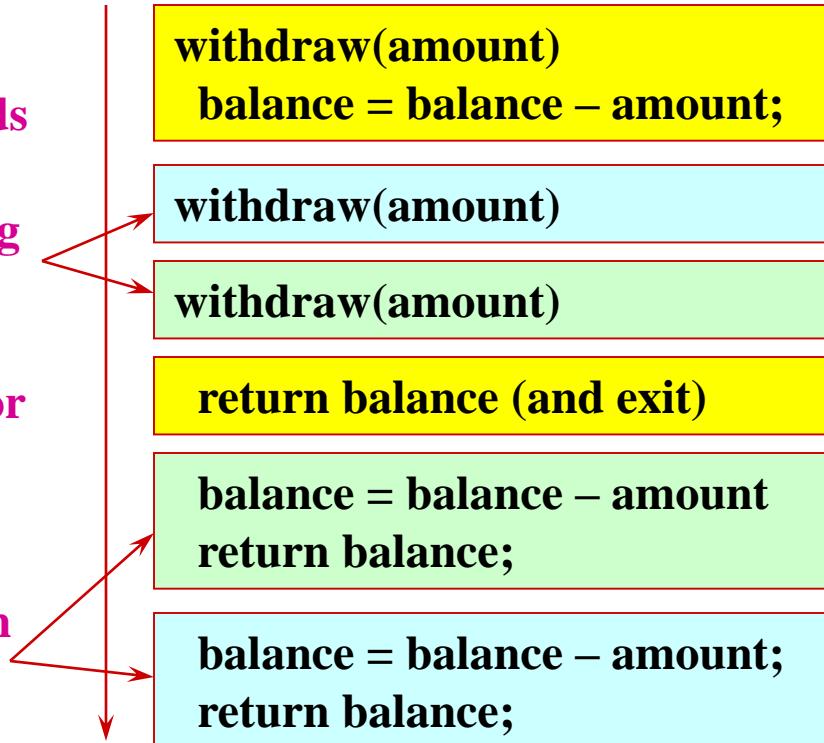


Account Example

```
Monitor account {  
    double balance;  
  
    double withdraw(amount) {  
        balance = balance - amount;  
        return balance;  
    }  
}
```

Threads
block
waiting
to get
into
monitor

When first thread exits, another can
enter. Which one is undefined.



- Hey, that was easy
- But what if a thread wants to wait inside the monitor?



Condition Variables

- `wait(condition, lock)`, `signal(condition)`
- Condition variables are NOT conditions
 - So don't EVER do this:
 - `if (condition variable){`
 - `...`
 - `}`
- Instead, it is a way for one thread to wait (if some resource is not available), and some other thread to wake it up once the resource becomes available
 - Sleep
 - Wake (also called as “signal”)
 - Wakeall (or “signalAll”)



Condition Variable & Lock

- Condition variable doesn't replace lock, instead it complements lock
- **wait(condition, lock)**
 - First release the lock, put the thread into the queue of the condition, if waking up, re-acquiring the lock
 - Once sleep returns, it is awaken by some other thread, and it also holds the lock
- **signal(condition)**
 - Wake up a thread waiting on the condition (queue)
- **broadcast(condition)**
 - Wake all the thread waiting on the condition (queue)



Condition Variables inside Monitors (1/2)

- A **condition variable inside a monitor** is associated with a condition needed for a thread to make progress once it is in the monitor.

```
Monitor M {  
    ... monitored variables  
    Condition c, d;
```

```
void enter_mon (...) {  
    if (extra property C not true) wait(c); // waits outside of the monitor's mutex  
    do what you have to do  
    if (extra property true D) signal(d); // brings in one thread waiting on d  
}
```



Condition Variables inside Monitors (2/2)

- Condition variables support three operations:
 - Wait – release monitor lock, wait for C/V to be signaled
 - So condition variables have wait queues, too
 - Signal – wakeup one waiting thread
 - Broadcast – wakeup all waiting threads
- Condition variables *are not* boolean objects
 - “if (condition_variable) then” ... does not make sense
 - “if (num_resources == 0) then wait(resources_available)” does
 - An example will make this clearer



Bounded Buffer Problem

- Problem: there is a set of resource buffers shared by producer and consumer threads
 - Producer inserts resources into the buffer set
 - Output, disk blocks, memory pages, processes, etc.
 - Consumer removes resources from the buffer set
 - Whatever is generated by the producer
- Producer and consumer execute at different rates
 - No serialization of one behind the other
 - Tasks are independent (easier to think about)
 - The buffer set allows each to run without explicit handoff
- Safety:
 - If nc is number consumed, np number produced, and N the size of the buffer, then $0 \leq np - nc \leq N$



Monitor Bounded Buffer

```
Monitor bounded_buffer {  
    Resource buffer[N];  
    // Variables for indexing buffer  
    // monitor invariant involves these vars  
    Condition not_full; // space in buffer  
    Condition not_empty; // value in buffer  
  
    void put_resource (Resource R) {  
        if (buffer array is full)  
            wait(not_full);  
        Add R to buffer array;  
        signal(not_empty);  
    }  
}
```

```
Resource get_resource() {  
    if (buffer array is empty)  
        wait(not_empty);  
    Get resource R from buffer array;  
    signal(not_full);  
    return R;  
}  
} // end monitor
```

- What happens if no threads are waiting when signal is called?
 - Signal is lost

Monitor Queues

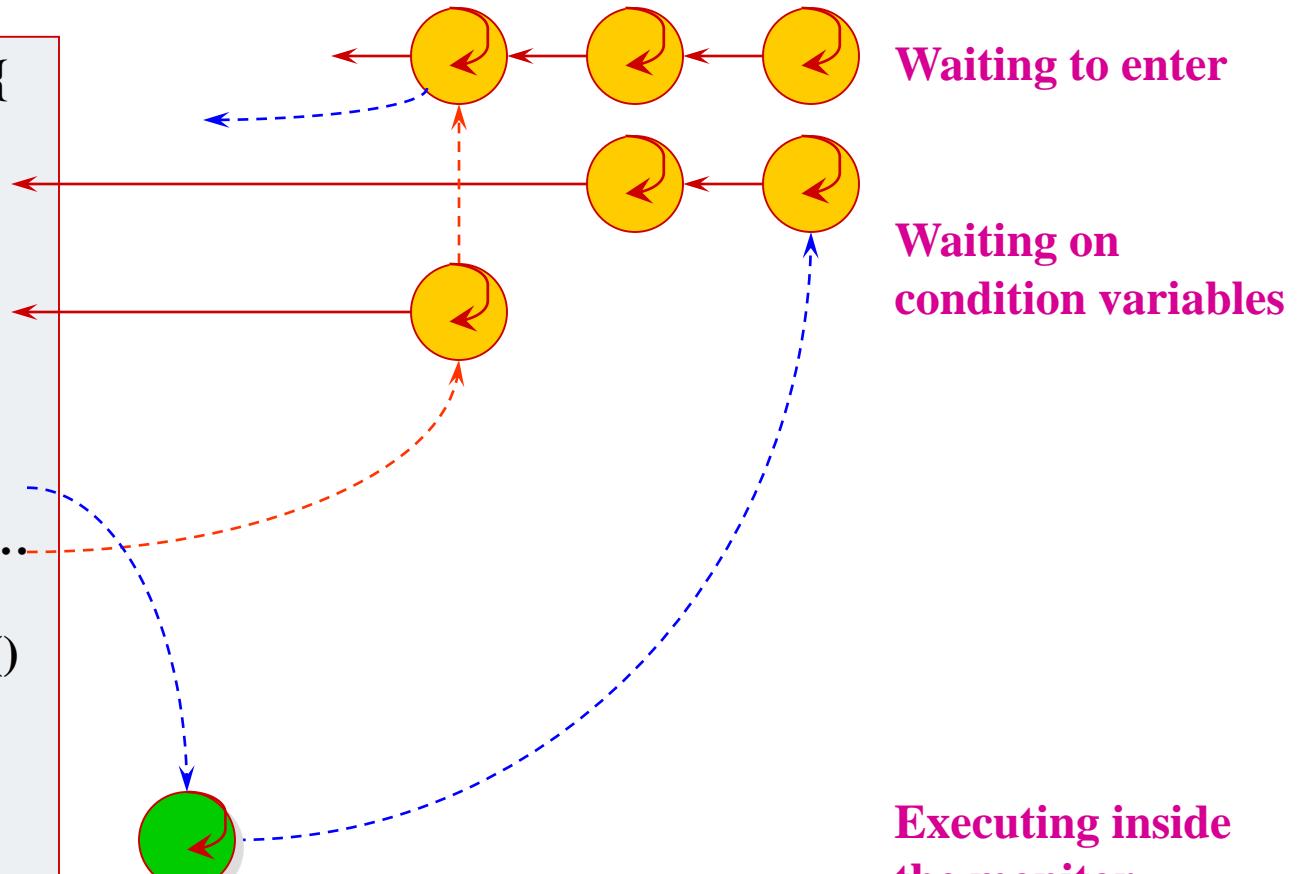
```

Monitor bounded_buffer {
    Condition not_full;
    ...other variables...
    Condition not_empty;

    void put_resource () {
        ...wait(not_full)...
        ...signal(not_empty)...
    }

    Resource get_resource ()
    {
        ...wait(not_empty)...
        ...signal(not_full)...
    }
}

```





Condition Vars != Semaphores

- Monitor with condition variables != semaphores
 - But they can implement each other
- Access to the monitor is controlled by a lock
 - `wait()` blocks the calling thread, and **gives up the lock**
 - To call wait, the thread has to be in the monitor (hence has lock)
 - `Semaphore::P()` just blocks the thread on the queue
 - `signal()` causes a waiting thread to wake up
 - **If there is no waiting thread, the signal is lost**
 - `Semaphore::V()` increases the semaphore count, allowing future entry even if no thread is waiting
 - Condition variables have no history



Signal Semantics

- There are two flavors of monitors that differ in the scheduling semantics of `signal()`
 - Hoare monitors (original)
 - `signal()` immediately switches from the caller to a waiting thread
 - The condition that the waiter was anticipating is guaranteed to hold when waiter executes
 - Mesa monitors (Mesa, Java)
 - `signal()` places a waiter on the ready queue, but signaler continues inside monitor
 - Condition is not necessarily true when waiter runs again
 - Returning from `wait()` is only a hint that something changed
 - Must recheck conditional case



Hoare vs. Mesa Monitors

- Hoare

```
if (empty)  
    wait(condition);
```

- Mesa

```
while (empty)  
    wait(condition);
```

- Tradeoffs

- Mesa monitors easier to use, more efficient
 - Fewer context switches, easy to support broadcast
- Hoare monitors leave less to chance
 - Easier to reason about the program



Monitor: Readers and Writers (1/2)

Using Mesa monitor semantics.

- Will have four methods: `StartRead`, `StartWrite`, `EndRead` and `EndWrite`
- Monitored data: `nr` (number of readers) and `nw` (number of writers) with the monitor invariant
- Two conditions:
 - `canRead`: $nw = 0$
 - `canWrite`: $(nr = 0) \wedge (nw = 0)$



Monitor: Readers and Writers (2/2)

```
Monitor RW {  
    int nr = 0, nw = 0;  
    Condition canRead, canWrite;  
  
    void StartRead () {  
        while (nw != 0) do wait(canRead);  
        nr++;  
    }  
  
    void EndRead () {  
        nr--;  
        if (nr==0) signal(canWrite)  
    }  
}
```

```
void StartWrite {  
    while (nr != 0 || nw != 0) do wait(canWrite);  
    nw++;  
}  
  
void EndWrite () {  
    nw--;  
    signal(canWrite);  
    broadcast(canRead);  
}  
} // end monitor
```

- Who protects accesses to nr and nw?
- Is there any priority between readers and writers? Any issue?
 - Starvation
- Why do we need “broadcast” in EndWrite?



Homework

1. Investigation Reports
 - Survey how locks are implemented in Linux (both for uniprocessor and multiprocessors)
 - Survey how semaphores are implemented in Linux (both for uniprocessor and multiprocessors)
 2. Exercise 2-57
 3. Exercise 2-60
 4. Exercise 2-62
-
- Present them in class
 - Student sign-up in WeChat group.
 - One student each question.



Summary

- Locks
 - Semaphores
 - Monitors and condition variables
-
- Next lecture: Discussions on Synchronization