



Operating Systems (A) (Honor Track)

Lecture 8: Memory Management

Yao Guo (郭耀)

Peking University

Fall 2021



Review: Threads

- The operating system as a large multithreaded program
 - Each process executes as a thread within the OS
- Multithreading is also very useful for applications
 - Efficient multithreading requires fast primitives
 - Processes are too heavyweight
- Solution is to separate threads from processes
 - Kernel-level threads much better, but still significant overhead
 - User-level threads even better, but not well integrated with OS



Review: Process vs. Thread

What are the main differences between processes and threads?



This Lecture

Memory Management

Overview

Memory Abstraction

Managing Free Memory

Buzz Words

Address space

**Physical/virtual
address**

Relocation

**Internal
fragmentation**

**External
fragmentation**

Compaction

Swapping


Memory is always limited

□ Parkinson's Law:

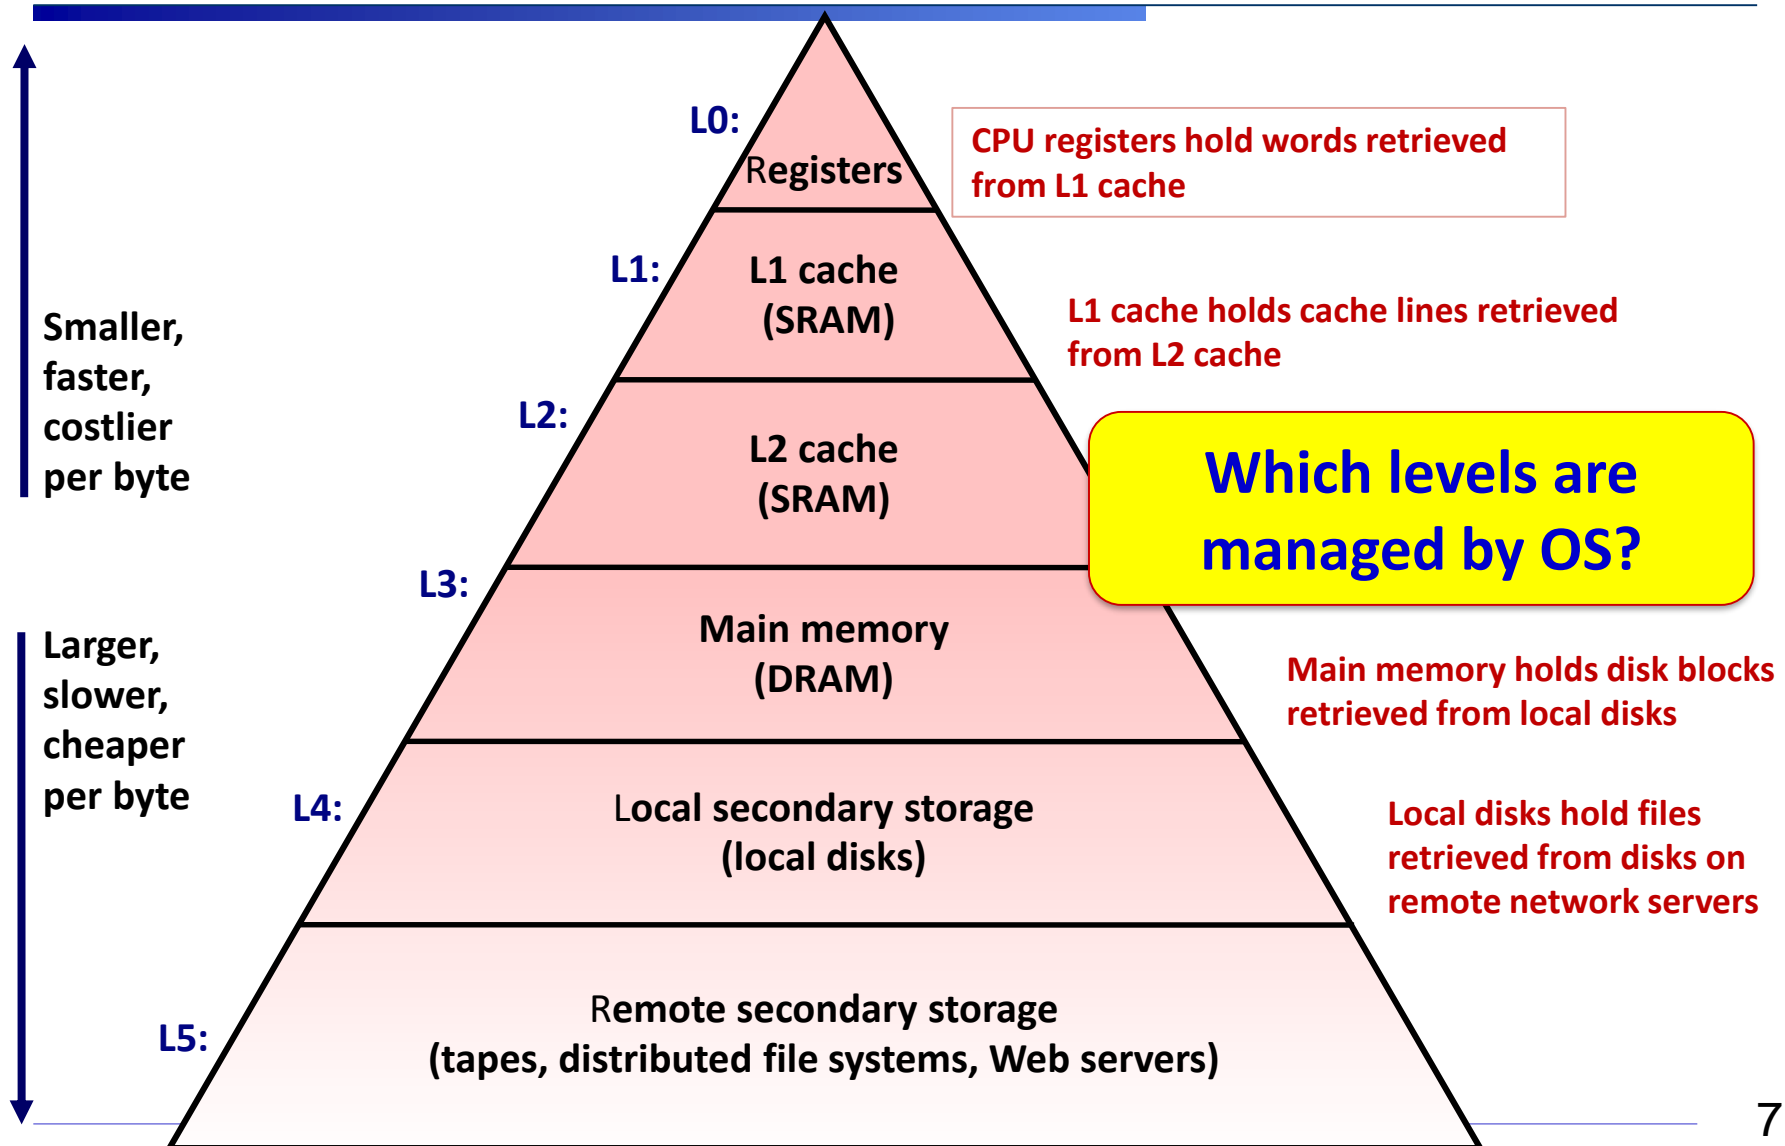
“Work expands so as to fill the time available for its completion”

□ For Memory:

“Programs expand to fill the memory available to hold them.”



Memory Hierarchy (from CSAPP)



Memory Manager

- **Memory manager**: the part of the operating system that manages (part of) the memory hierarchy
- The job of a memory manager is to **efficiently** manage memory:
 - **Keep track of** which parts of memory are in use,
 - **Allocate** memory to processes when they need it, and
 - **Deallocate** it when they are done



Memory Management

- **Goals** of memory management
 - To provide a convenient abstraction for programming
 - To allocate scarce memory resources among competing processes to maximize performance with minimal overhead
- **Mechanisms**
 - Physical and virtual/logical addressing
 - Techniques: partitioning, paging, segmentation
 - Page table management, TLBs, VM tricks
- **Policies**
 - Page replacement algorithms

Virtual Memory

- The abstraction that the OS will provide for managing memory is **virtual memory (VM)**
 - Virtual memory enables a program to execute with less than its complete data in physical memory
 - A program can run on a machine with less memory than it “needs”
 - Can also run on a machine with “too much” physical memory
 - Many programs do not need all of their code and data at once (or ever) – no need to allocate memory for it
 - OS will adjust amount of memory allocated to a process based upon its behavior
 - VM requires hardware support and OS management algorithms to pull it off



This Lecture

Memory Management

Overview

Memory Abstraction

Managing Free Memory

When there were none...

- The simplest memory abstraction is to have **no abstraction** at all.
 - early mainframe computers (before 1960),
 - early minicomputers (before 1970), and
 - early personal computers (before 1980)
- Every program simply saw the physical memory.
- Can we run two programs at the same time?
 - How about two processes?
 - And two threads?



The Need for Memory Abstraction


- Exposing physical memory to processes is bad
 - It can easily trash the operating system, even when there is only one process
- It is also difficult to support multiprogramming
 - Want **multiple processes in memory** at once
 - Overlap I/O and CPU of multiple jobs
 - **Need protection** – restrict which addresses jobs can use
 - **Fast translation** – lookups need to be fast
 - **Fast change** – updating memory hardware on context switch

Address Space

- An (virtual) **address space** is the set of addresses that a process can use to address memory.
- Each process has its own address space, independent of those belonging to other processes
 - With some exceptions, such as **shared memory**
- **Relocation**: logical address \Rightarrow physical address
 - Static relocation, or
 - Dynamic relocation


Static Relocation

- **Static relocation**: convert to physical addresses for all memory accesses during loading
 - Can be implemented in software
 - High overhead



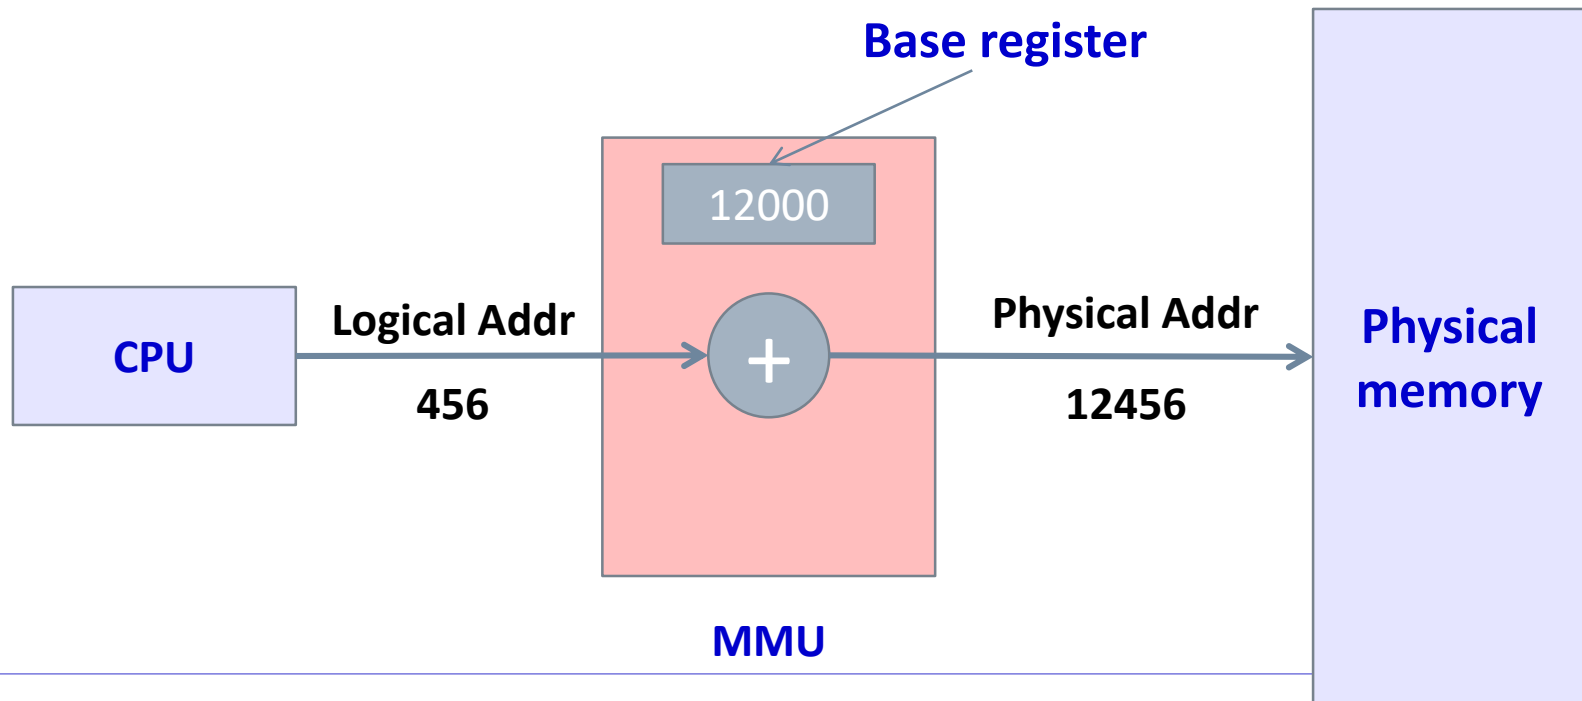
Dynamic Relocation

- **Dynamic relocation**: convert to physical addresses when an instruction is executed
 - Hardware support?



Dynamic Relocation Implementation

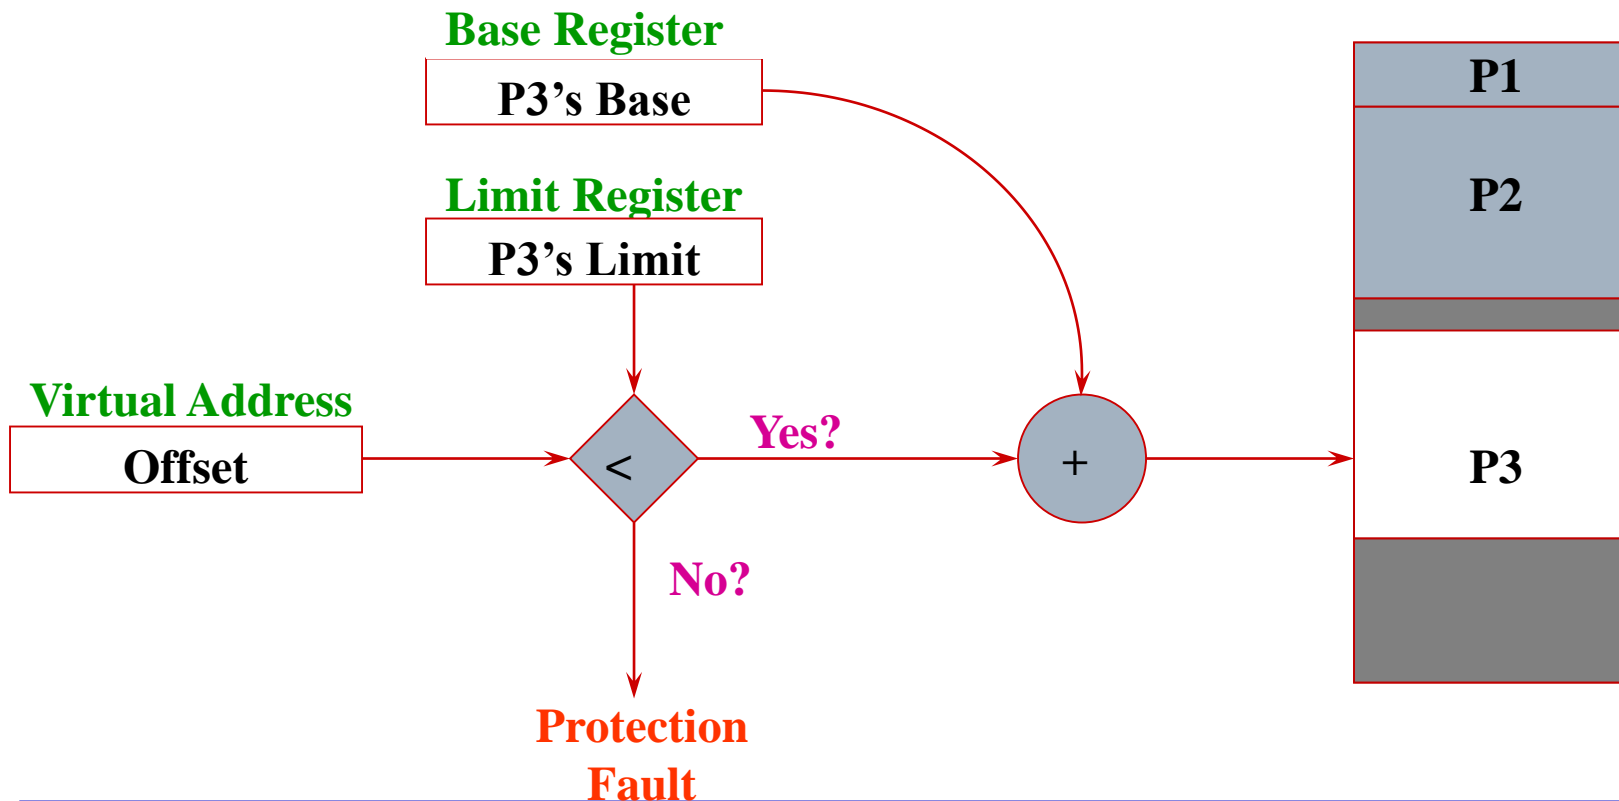
- **MMU**: Memory Management Unit
- Two special hardware registers: the **base** and **limit** registers.
 - Why do we need the limit register?



MMU Implementation

□ Memory protection

- If (physical address > base + limit) then exception fault



When Memory is Full...

- **Swapping:** swap in/out a process from/to disk

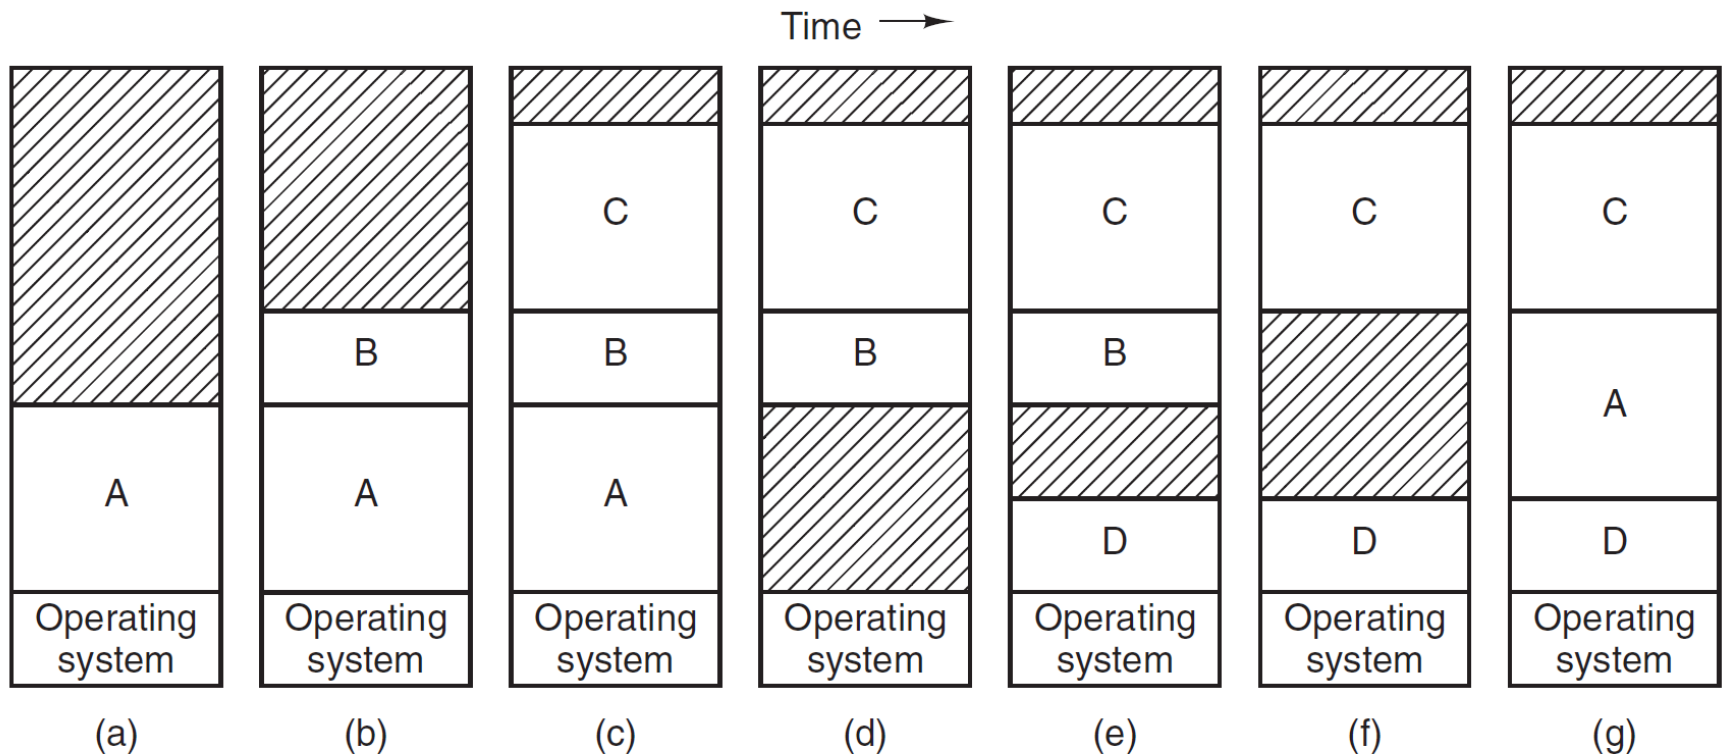


Figure 3-4. Memory allocation changes as processes come into memory and leave it. The shaded regions are unused memory.



This Lecture

Memory Management

Overview

Memory Abstraction

Managing Free Memory

Managing Free Memory

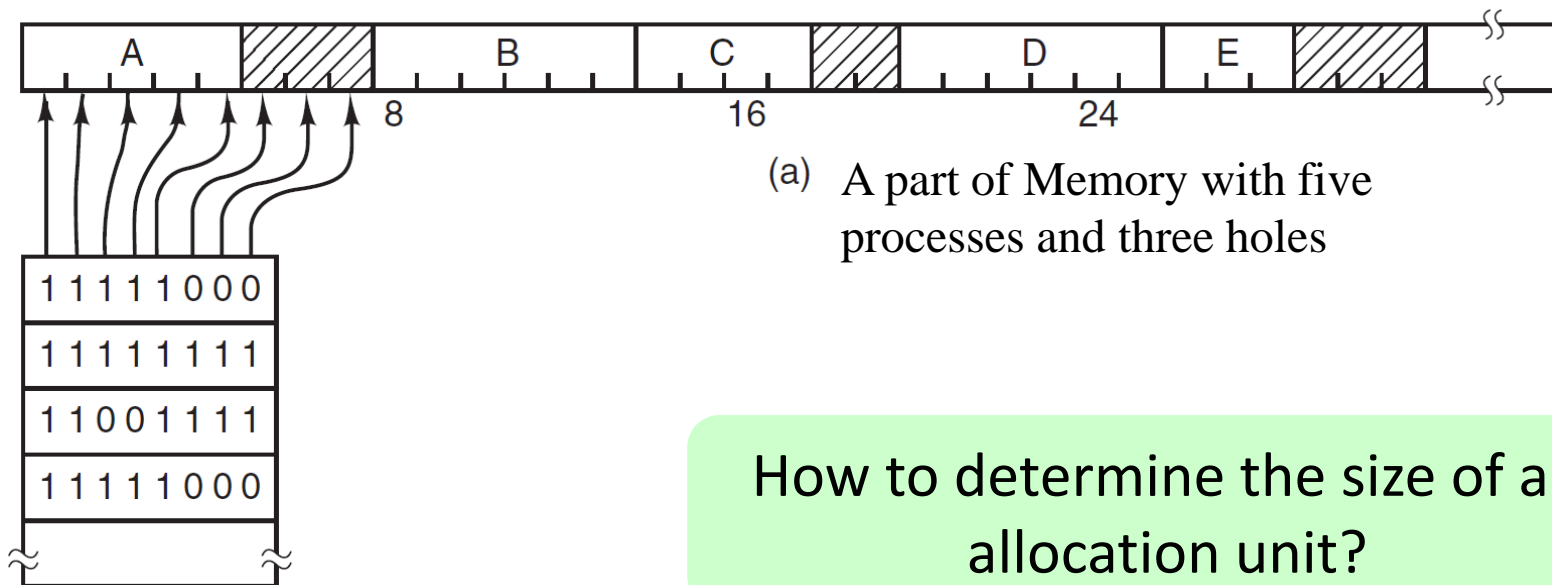
- OS is responsible of managing dynamically allocated memory

- Methods:
 - Bitmaps
 - Free lists

- For Linux
 - Buddy system
 - Slab allocator

Memory Management with Bitmaps


- Each **allocation unit** is represented with 1 bit in the bitmap



How to determine the size of an allocation unit?

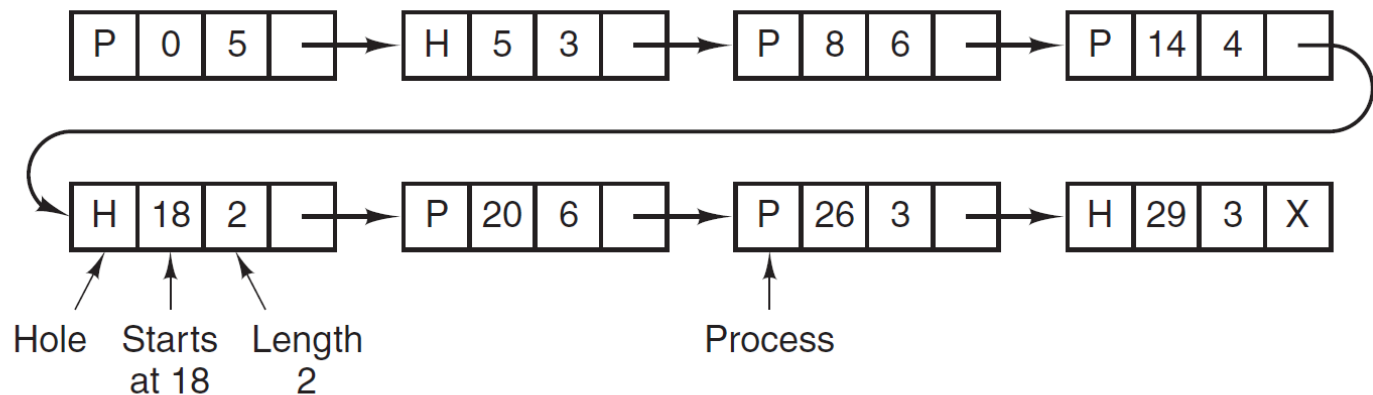
Memory Management with Linked Lists

- A linked list of allocated and free memory segments,
 - where a segment either contains a process or is an empty hole between two processes.



(a) A part of Memory

Linked lists
with allocated
and free units



(c)

Memory Allocation Algorithms

- How to allocate memory for a created process?
- First fit
 - Find the first hole that is big enough
 - Problem: creates average size holes
- Next fit
 - starts searching the list from the place where it left off last time
 - (Is it better than first fit?)
- Best fit
 - find the smallest hole that is adequate
 - Problem: creates small holes that can't be used

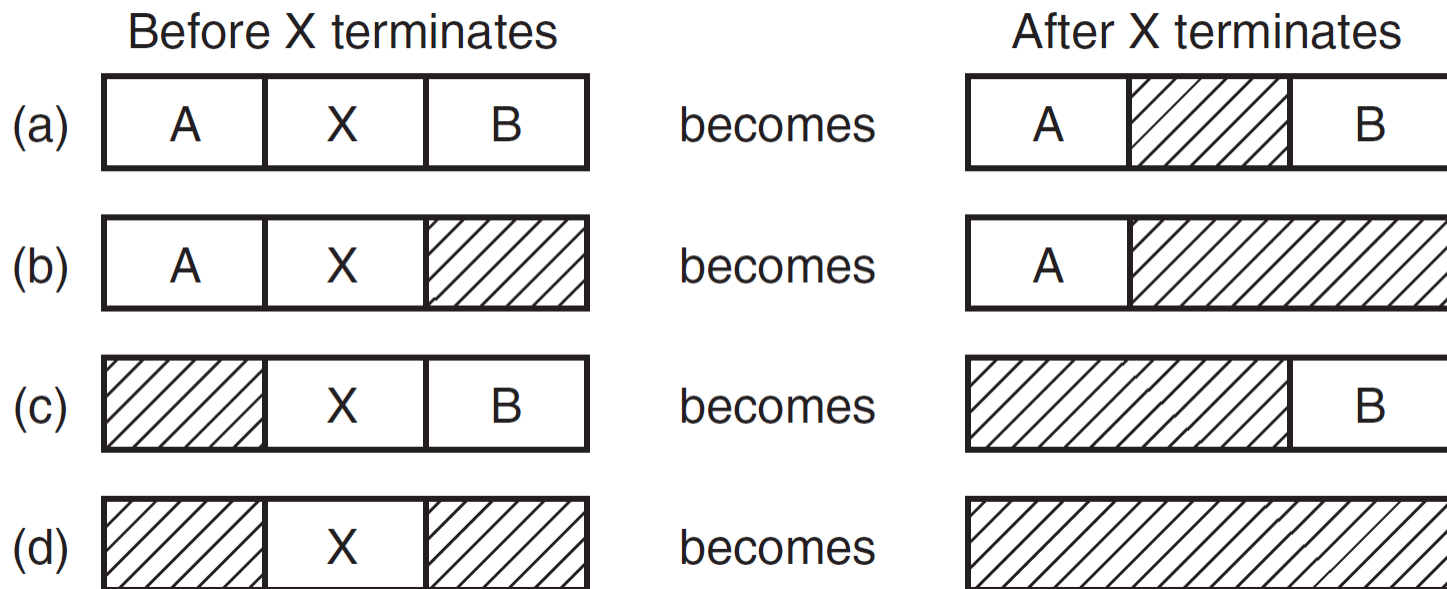


Memory Allocation Algorithms (cont.)

- How to allocate memory for a created process?
- Worst Fit
 - Use the largest available hole
 - Problem: gets rid of large holes making it difficult to run large programs
- Quick Fit
 - Maintains separate lists for some of the more common sizes requested
 - When a request comes for placement it finds the closest fit
 - Very fast, but merging is expensive

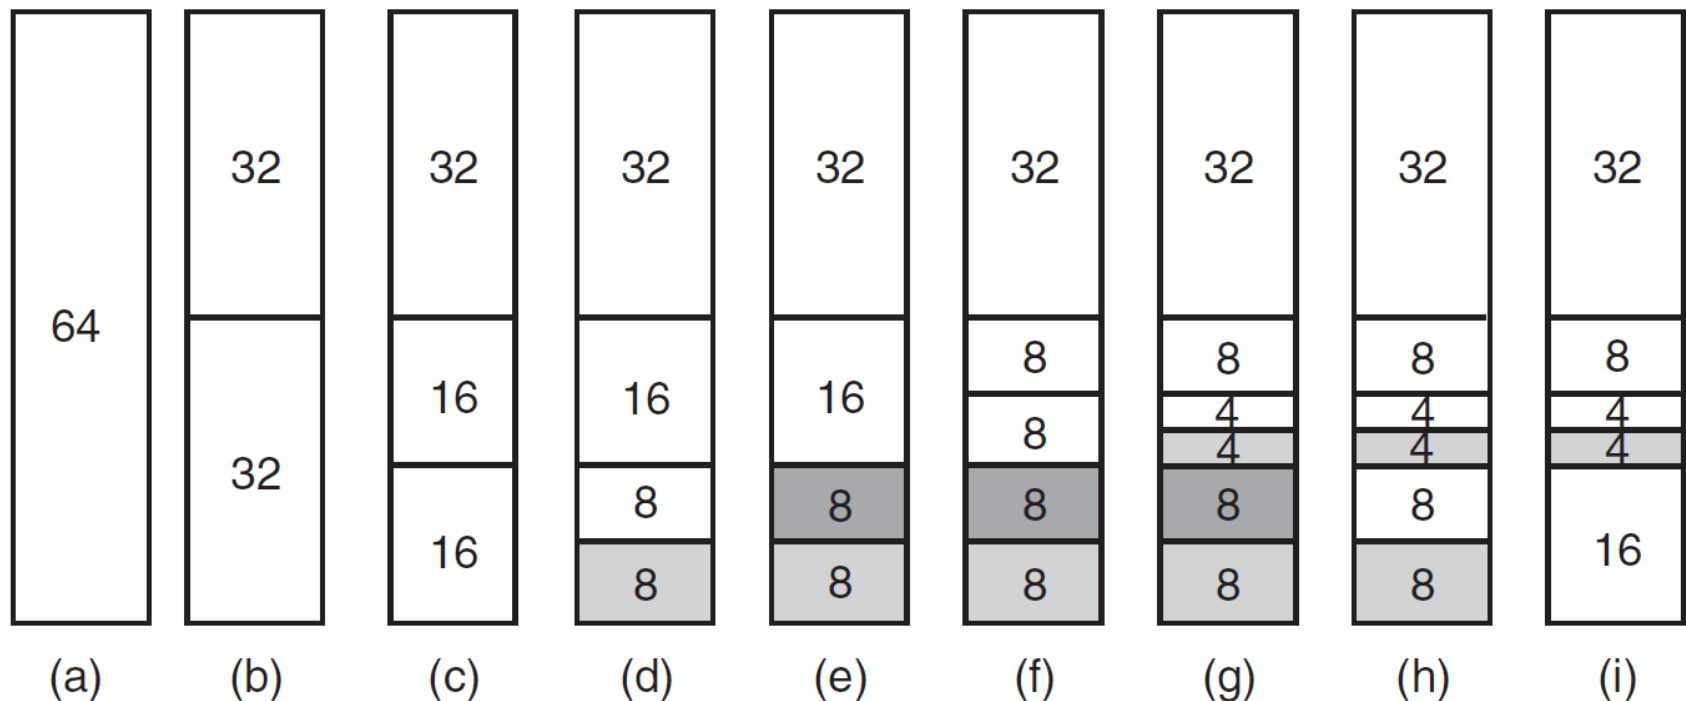
Memory Compaction

- A terminating process normally has two neighbors (except when it is at the very top or bottom of memory)
- Merge it with neighbors when it is terminated



Buddy System (Linux)

- It is also called the **buddy algorithm**
 - All free segments are in size of power of 2 (2^n)
 - Find the smallest hole larger than the requested size



.....and other allocators

- Benefit: fast allocation possible
 - All blocks with the same size are put in a list
 - With an array pointing to each list with different sizes
- Problem: internal fragmentation is high, because the smallest unit is a page.
- Solution: **Slab allocator**
 - Managing small chunks (bytes within a page) separately
 - Used in **kmalloc**: allocates physically contiguous memory regions in the kernel address space
 - There is also **vmalloc**: allocates logically contiguous memory regions (with some performance degradation)



When there is not enough memory

- What can an OS do when it does not have enough memory?
 - Memory compaction
 - Memory overlaying
 - Swapping
 - Virtual memory (next lecture)


Memory Overlaying

- Static memory allocation
 - Different functions can reuse the same memory space if they are not dependent on each other
 - Load each overlaid module when it is called
 - Needs to know the calling relationship of all functions (subroutines)

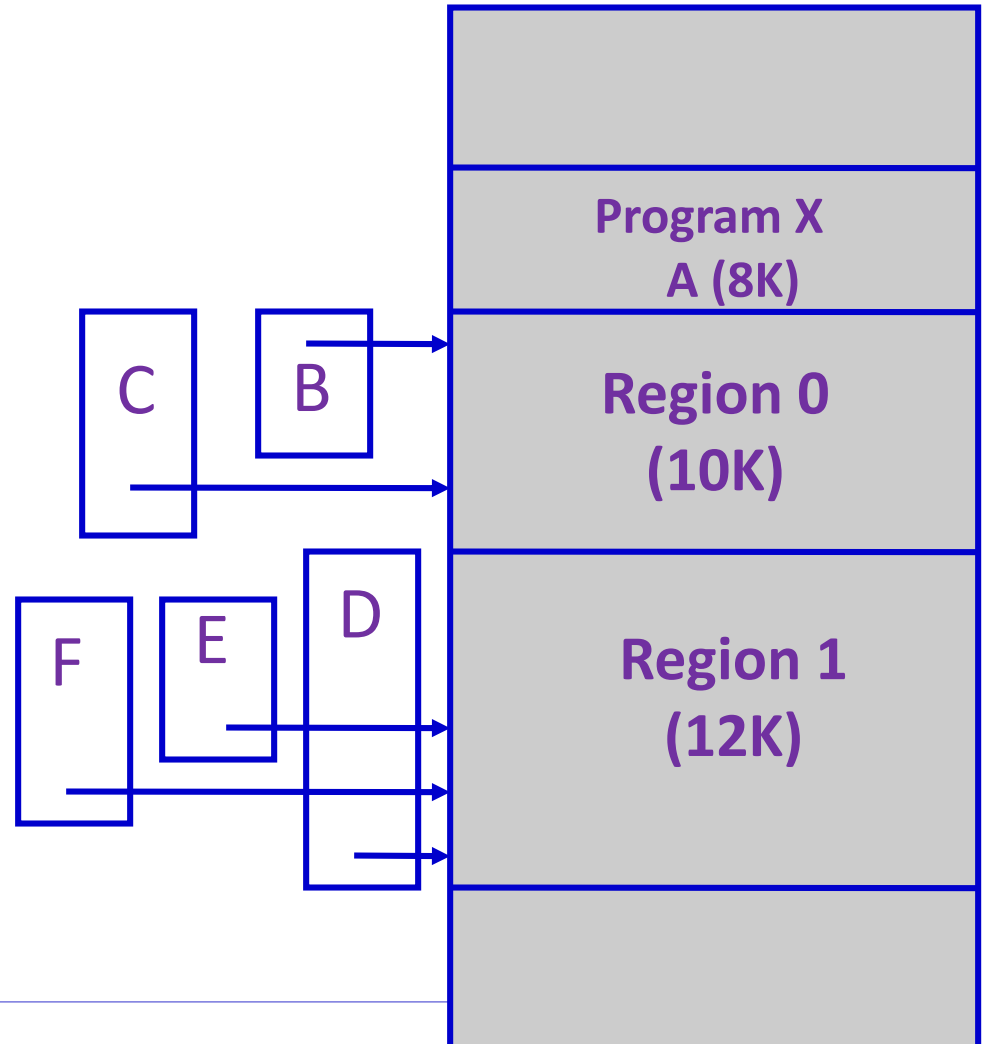
- Used in early systems
 - Such as early FORTRAN

Example of Overlaying

Memory requirements:
30K



Calling graph of Program X:
Total requirements: 52K



Summary

- Overview of Memory Management
 - Memory abstraction: address space
 - OS + hardware translate virtual address into physical addresses
- Static and dynamic relocation
- Memory allocation algorithms
 - First fit, second fit, best fit, worst fit, quick fit
- Managing free memory
 - Bitmap, linked lists, buddy, slab, (and [slub](#))
- Next lecture: Virtual Memory