



Operating Systems (A) (Honor Track)

Lecture 11: Windows Virtual Memory

Yao Guo (郭耀)

Peking University

Fall 2021

Acknowledgements: Prof. Xiangqun Chen & Tao Wang at PKU and Prof. Yuanyuan Zhou at UCSD

Buzz Words

Protected Mode

**Virtual Address
Descriptor (VAD)**

Working set

Self-Map



This Lecture

Windows MM

Intel x86 Virtual Memory

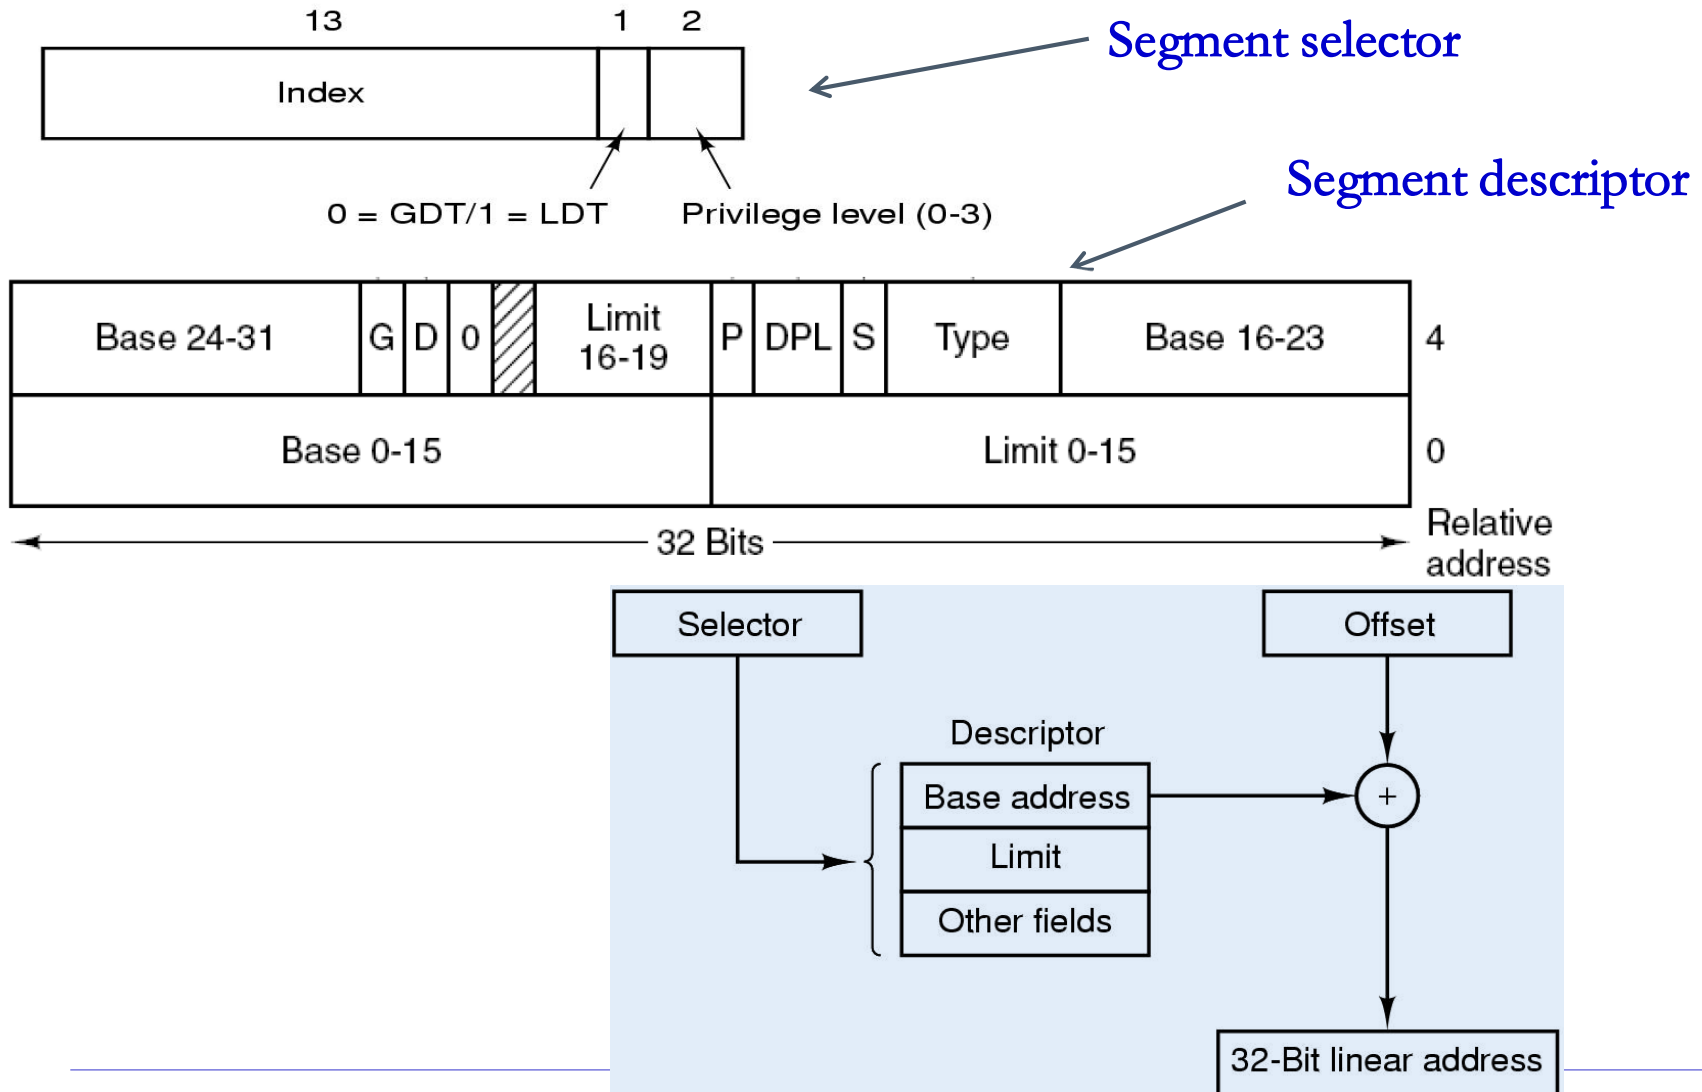
Windows MM




Address Mode for x86

- Real mode (8086)
 - Using physical address
 - For DOS
- Protected mode (80286)
 - Protected virtual address mode
 - Supports virtual memory, paging and safe multi-tasking
 - Supports both 16 and 32 bits
- Intel 8086 virtual mode (80386)
 - Hardware virtualization
 - Run real mode in the protected mode


Protected Mode



Logical to Physical Address



(a)



i386 Page Table Entries

PDE (Page Directory Entry)

PFN	Avail	G	P S	0	A	P C D	P W T	U /S	R / W	P
-----	-------	---	--------	---	---	-------------	-------------	---------	-------------	---

PTE (Page Table Entry)

PFN	Avail	G	0	D	A	P C D	P W T	U /S	R / W	P
-----	-------	---	---	---	---	-------------	-------------	---------	-------------	---

63	62	52	51	12	11	9	8	7	6	5	4	3	2	1	0
N X	AVL	Physical page number			AVL	G	P A T	D	A	P C D	P W T	U / S	R / W	P	

NX No eXecute

AVL AVaiLable to the OS

G Global page

PAT Page Attribute Table

D Dirty (modified)

A Accessed (referenced)

PCD Page Cache Disable

PWT Page Write-Through

U/S User/Supervisor

R/W Read/Write access

P Present (valid)

**Intel x86 or
AMD x64 PTEs**



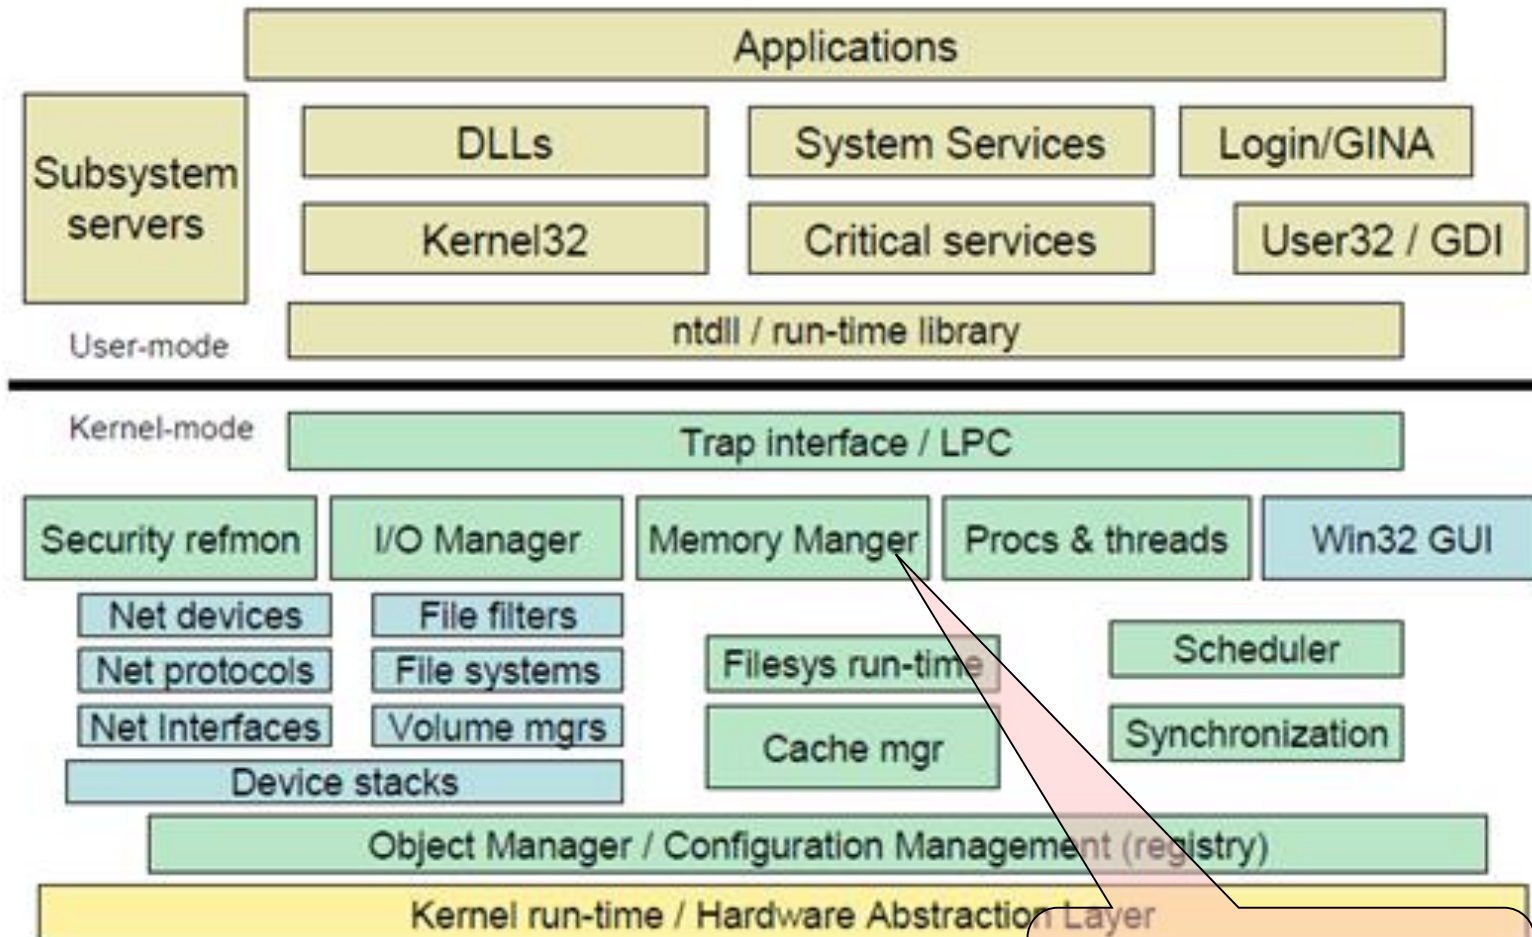
This Lecture

Windows MM

Intel x86 Virtual Memory

Windows MM

The Windows Architecture



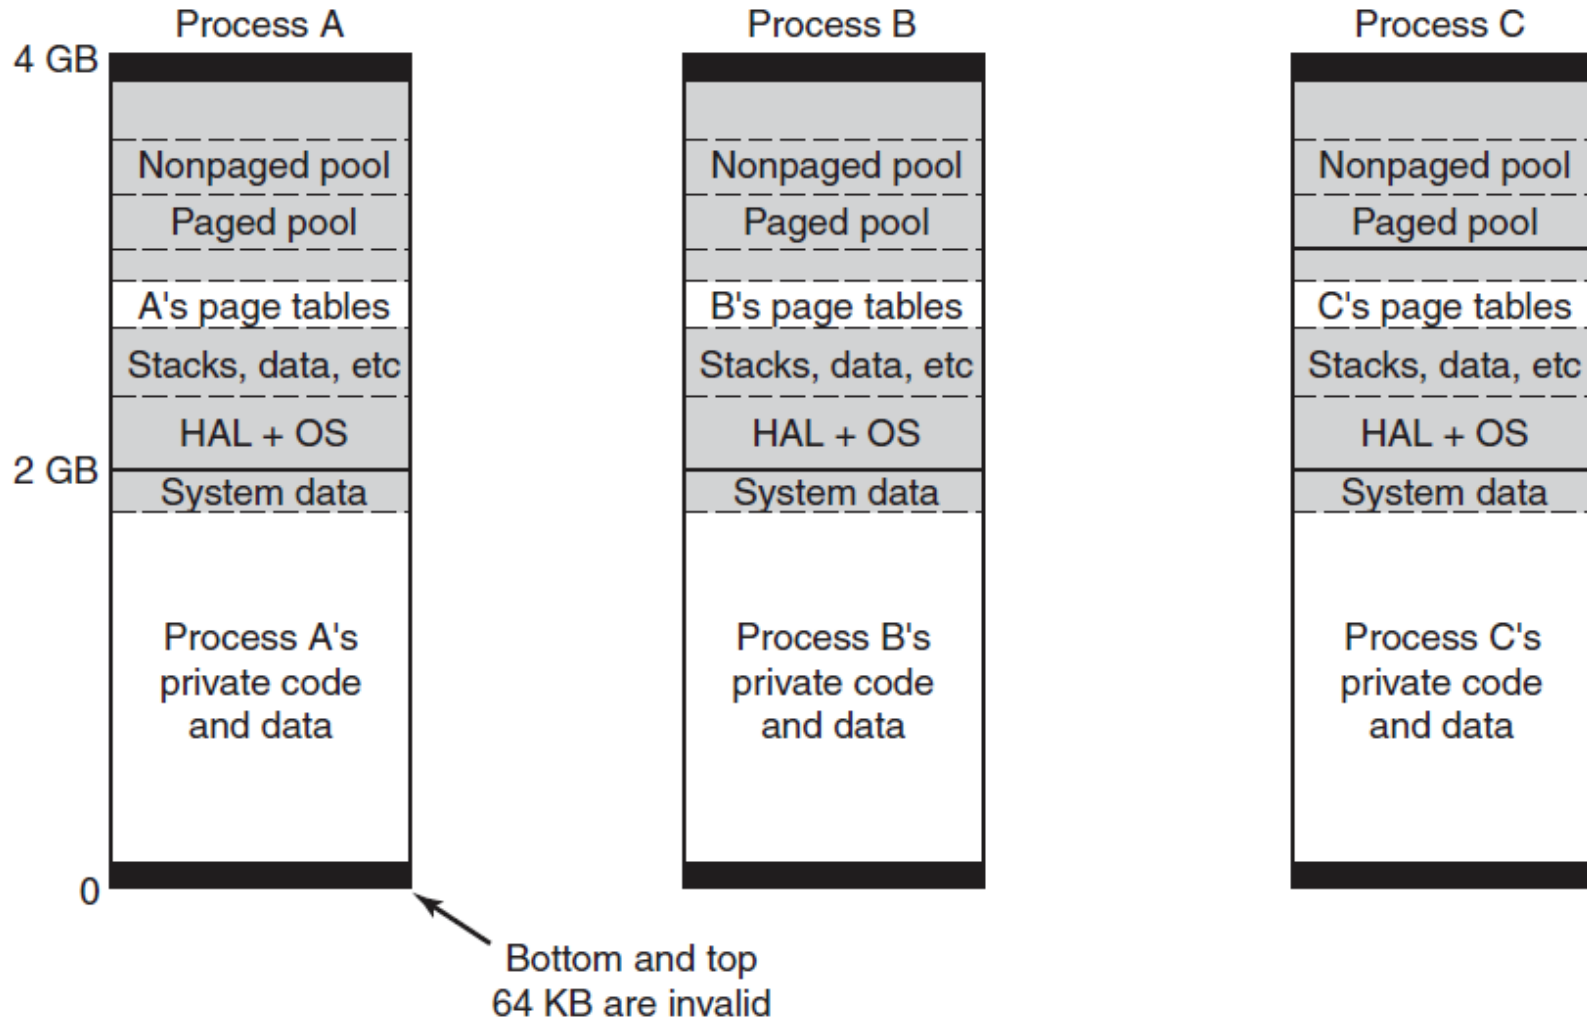
MM:
Ntoskrnl.exe



Windows Virtual Memory

- x86
 - Virtual addresses: 32 bits
 - Address space for each process: 4GB
 - Kernel: 2GB
 - User : 2GB
- x86-64
 - Large enough
- Demand paging
 - Page size: 4KB
 - (2MB pages are also used)

Virtual Address Space Layout





Virtual Address Allocation

- Each page of virtual addresses can be in one of three states:
 - **Invalid**: not currently mapped to a memory section object and a reference to it causes a page fault
 - **Committed**: code or data is mapped onto a virtual page
 - **Reserved**: invalid but has the property that those virtual addresses will never be allocated by the memory manager for another purpose
 - function as **guard pages** to keep the stack from growing too far and overwriting other process data.



Windows MM System Calls

Win32 API function	Description
VirtualAlloc	Reserve or commit a region
VirtualFree	Release or decommit a region
VirtualProtect	Change the read/write/execute protection on a region
VirtualQuery	Inquire about the status of a region
VirtualLock	Make a region memory resident (i.e., disable paging for it)
VirtualUnlock	Make a region pageable in the usual way
CreateFileMapping	Create a file-mapping object and (optionally) assign it a name
MapViewOfFile	Map (part of) a file into the address space
UnmapViewOfFile	Remove a mapped file from the address space
OpenFileMapping	Open a previously created file-mapping object

Windows MM Implementation

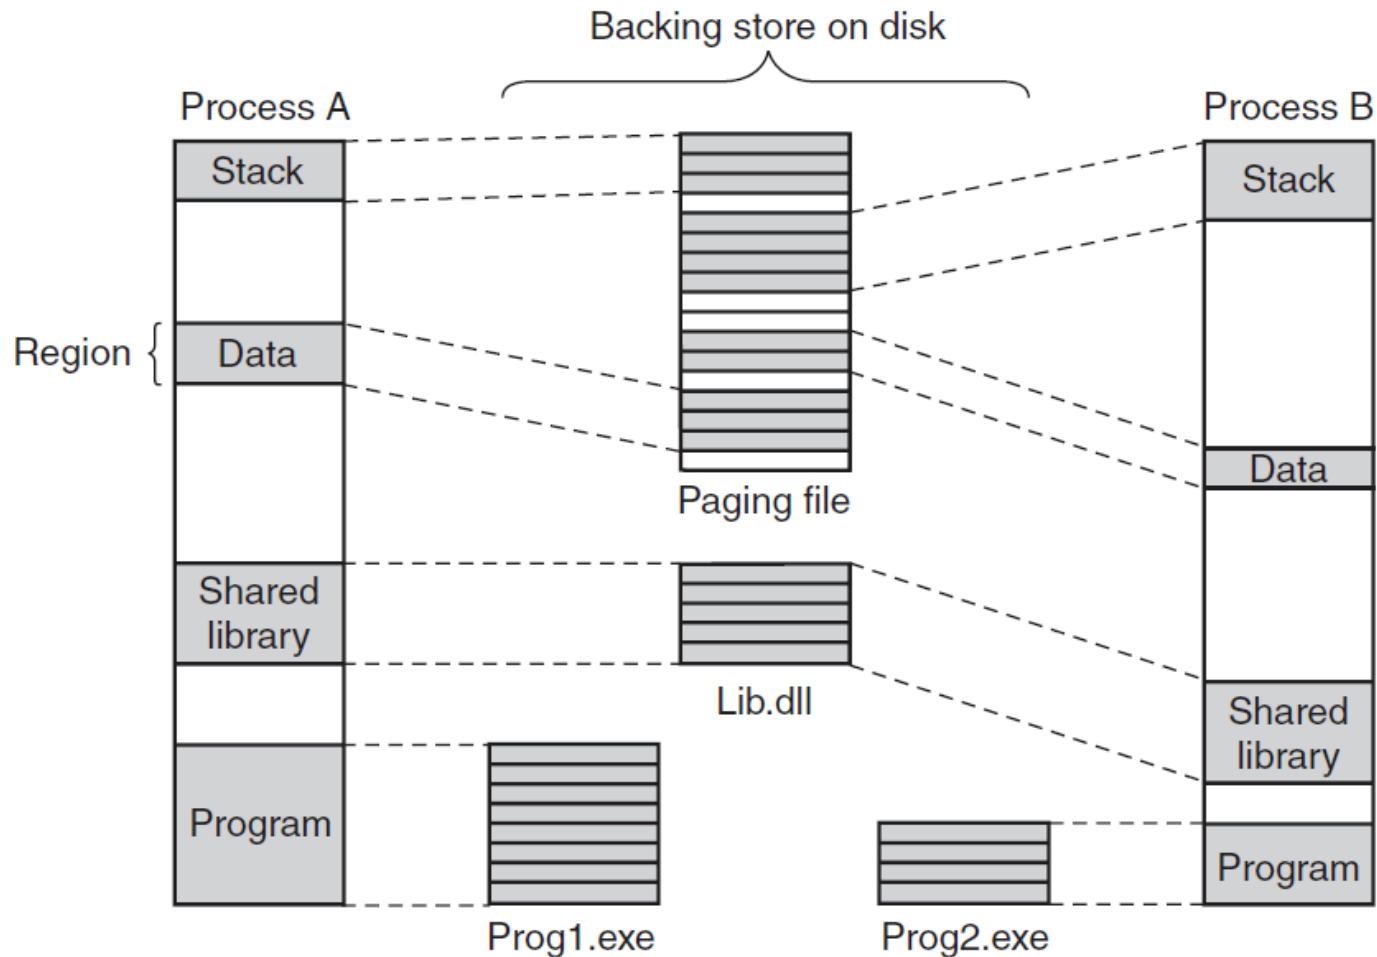


Figure 11-30. Mapped regions with their shadow pages on disk. The *lib.dll* file is mapped into two address spaces at the same time.



VAD (Virtual Address Descriptor)

- MM creates a VAD (Virtual Address Descriptor) for each process
 - Listing the range of addresses mapped, the section representing the backing store file and offset where it is mapped, and the permissions
 - An address space is completely defined by the list of its VADs
 - VADs are organized into a balanced tree

Page Fault Handling

- Types of pages faults:
 - The page referenced is not committed.
 - Access to a page has been attempted in violation of the permissions.
 - A shared copy-on-write page was about to be modified.
 - The stack needs to grow.
 - The page referenced is committed but not currently mapped in.


- How does each of these types occur?

Hard Faults vs. Soft Faults

- **Hard fault:** needs reading from disk
- **Soft fault:** the memory manager can satisfy a page fault by finding the needed page in memory rather than reading it in from disk
- **When might a soft fault occur?**
 - A shared page has already been mapped into another process
 - Only a new zero page is needed
 - The needed page was trimmed from the process' working set but is being requested again before it has had a chance to be reused
 - Pages have been compressed to effectively increase the size of physical memory.

Self-Map Entries

- The **self-map entries** are used to map the physical pages of the page tables and page directory into kernel virtual addresses (shown for 32-bit PTEs).



Self-map: `PD[0xc0300000 >> 22]` is PD (page-directory)

Virtual address (a): `(PTE *) (0xc0300c00)` points to `PD[0x300]` which is the self-map page directory entry

Virtual address (b): `(PTE *) (0xc0390c84)` points to PTE for virtual address `0xe4321000`

Self-Map Entries

MiGetPdeAddress():

Given a virtual address va , compute its PDE

$((\text{PMMPTE})((((\text{ULONG})(va)) \gg 22) \ll 2) + \text{PDE_BASE}))$

MiGetPteAddress():

Given a virtual address va , compute its PTE

$((\text{PMMPTE})((((\text{ULONG})(va)) \gg 12) \ll 2) + \text{PTE_BASE}))$



Page Replacement Algorithm

□ Working set

- Each process' working set is described by two parameters: the minimum size and the maximum size.

- Default minimum: 20–50 pages
- Default maximum: 45–345 pages

□ The working set manager runs every second

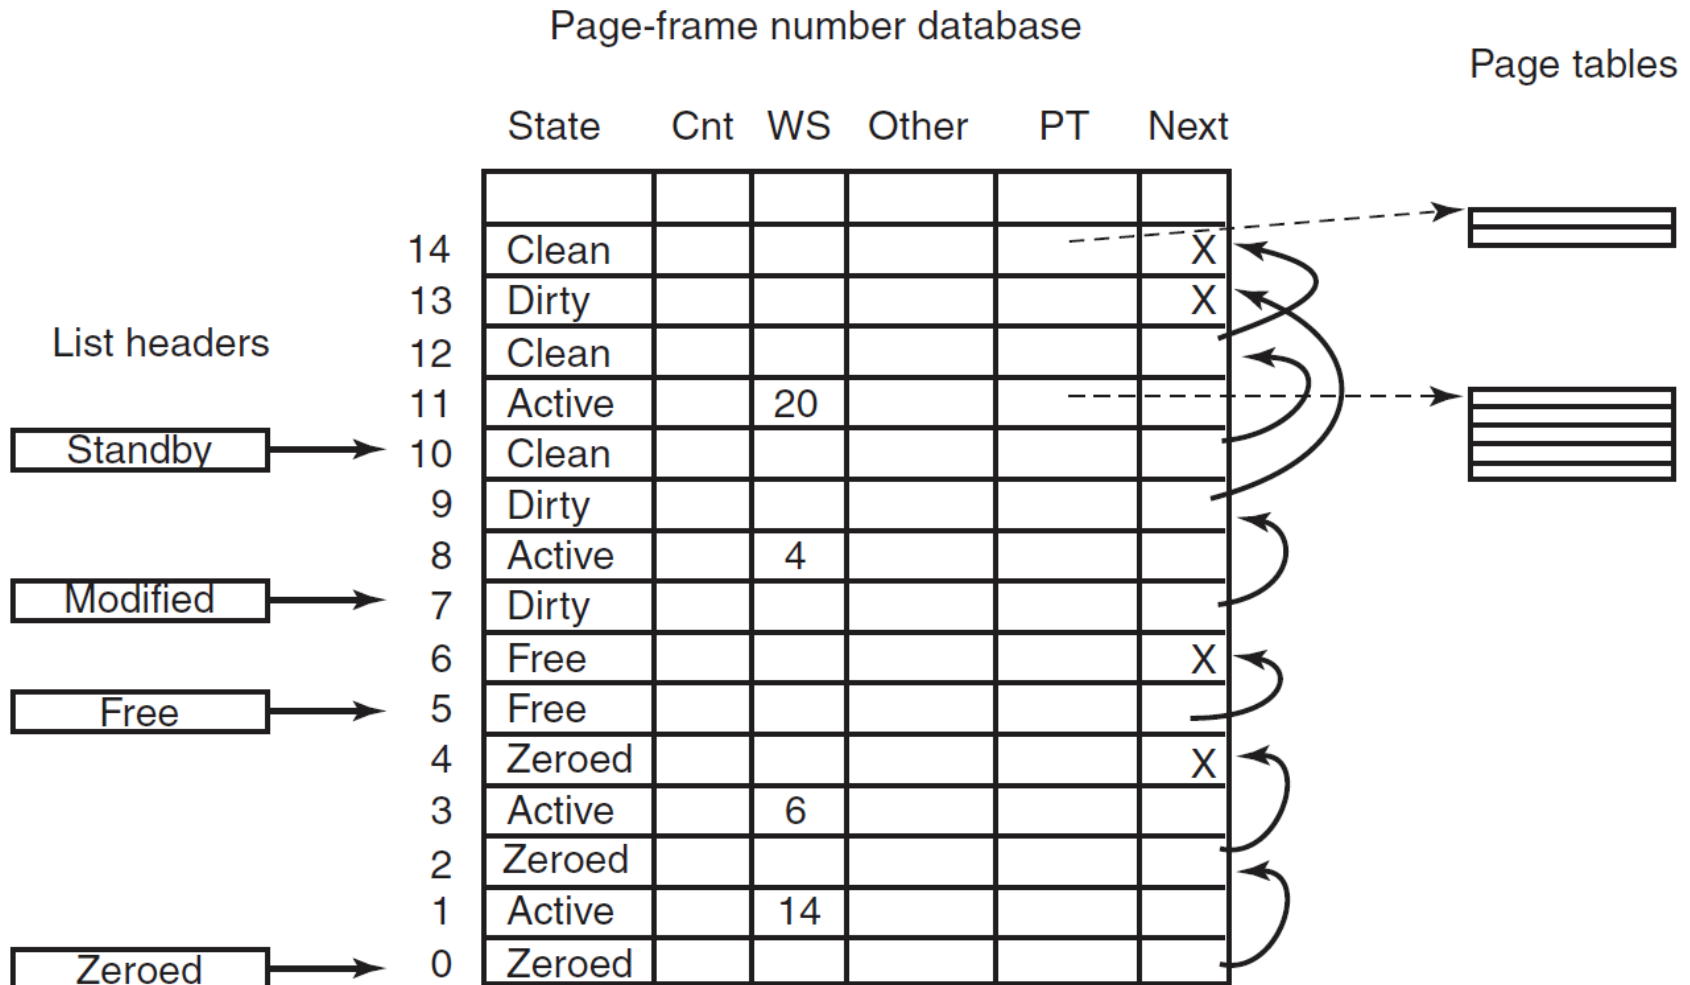
- The working-set manager throttles the amount of work it does to keep from overloading the system.

Physical Memory Management

- PFN (Page Frame Number) Database
 - All pages in the system either are referenced by a valid page-table entry or are on one of these five lists

- Different page frame lists
 - The free list
 - The standby list
 - The modified list
 - The zeroed list (free & zeroed)
 - The bad memory page list (frames with hardware errors)

Physical Memory Management



Transitions between Page Lists

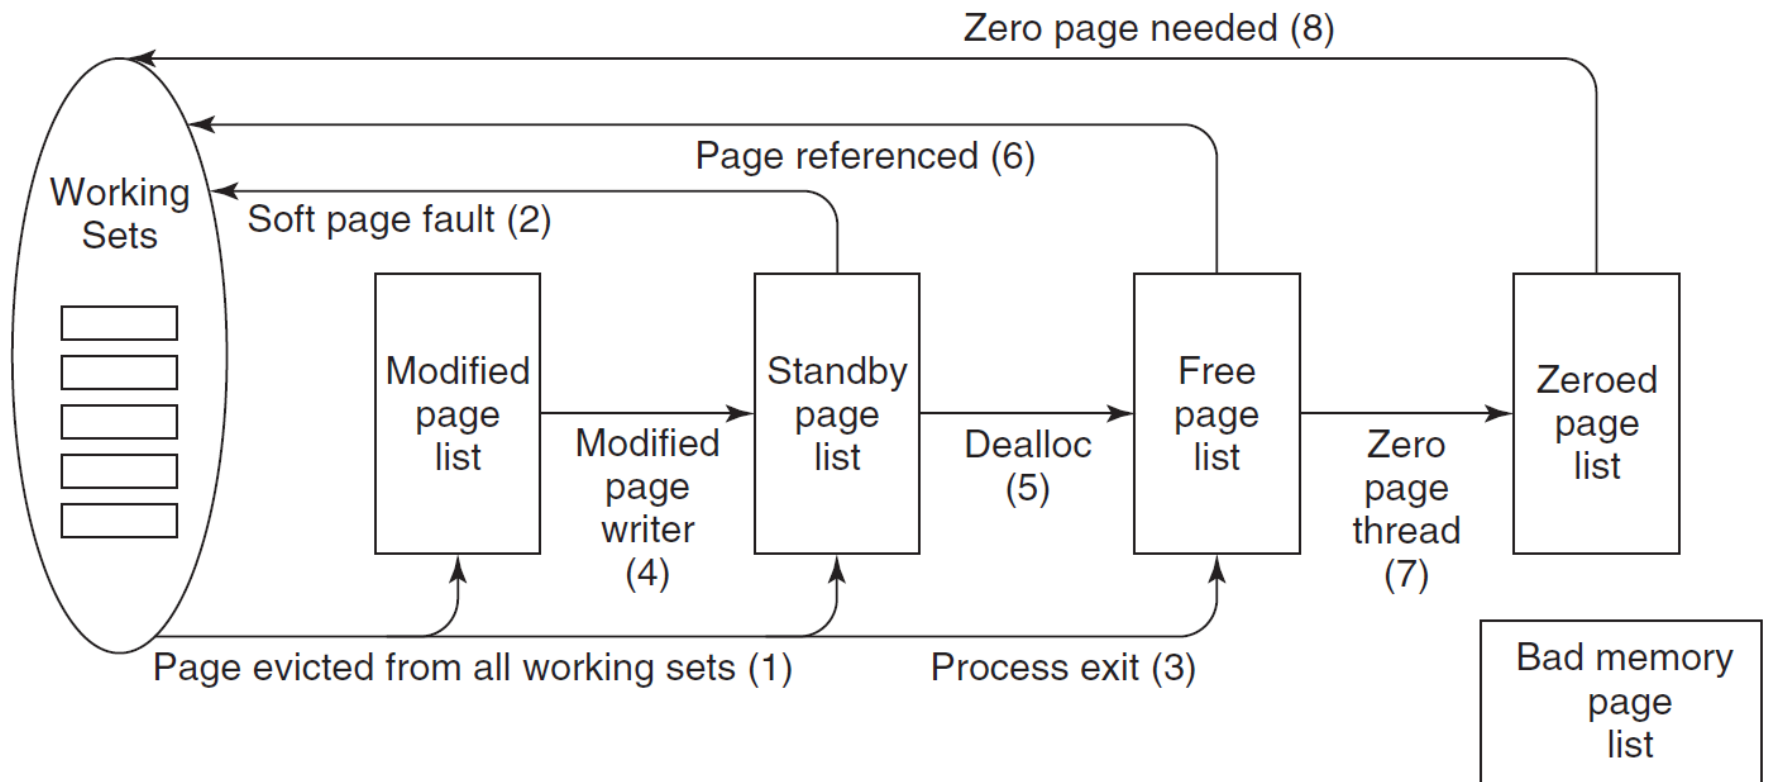


Figure 11-34. The various page lists and the transitions between them.

Summary

- Windows Memory Management
 - Virtual Address Descriptor
 - Self-Map Entries
 - Various Kinds of Page Lists

- Next Lecture: Scheduling