



Operating Systems (A) (Honor Track)

Lecture 19: File System Examples

Yao Guo (郭耀)

Peking University

Fall 2021



This Lecture

File System Examples

Fast File System (FFS)

Journaling File Systems

Log-structure File Systems

FAT

GFS



Fast File System (1980s)

- ❑ Old Unix FS: performance degradation over time
- ❑ FFS: First disk-aware file system
 - Bitmaps
 - Locality groups
 - Rotated superblocks
 - Large blocks
 - Fragments
 - Smart allocation policy
- ❑ FFS inspired modern files systems, including ext2 and ext3
- ❑ FFS also introduced several new features:
 - long file names
 - atomic rename
 - symbolic links



The Linux File System

- MINIX
 - file names of 14 characters and file size of 64 MB
- ext
 - allowed file names of 255 characters and files of 2 GB
 - Performance is worse than MINIX
- ext2
 - long file names, long files, and better performance
- ext3
- ext4
- btrfs, xfs, zfs, etc.

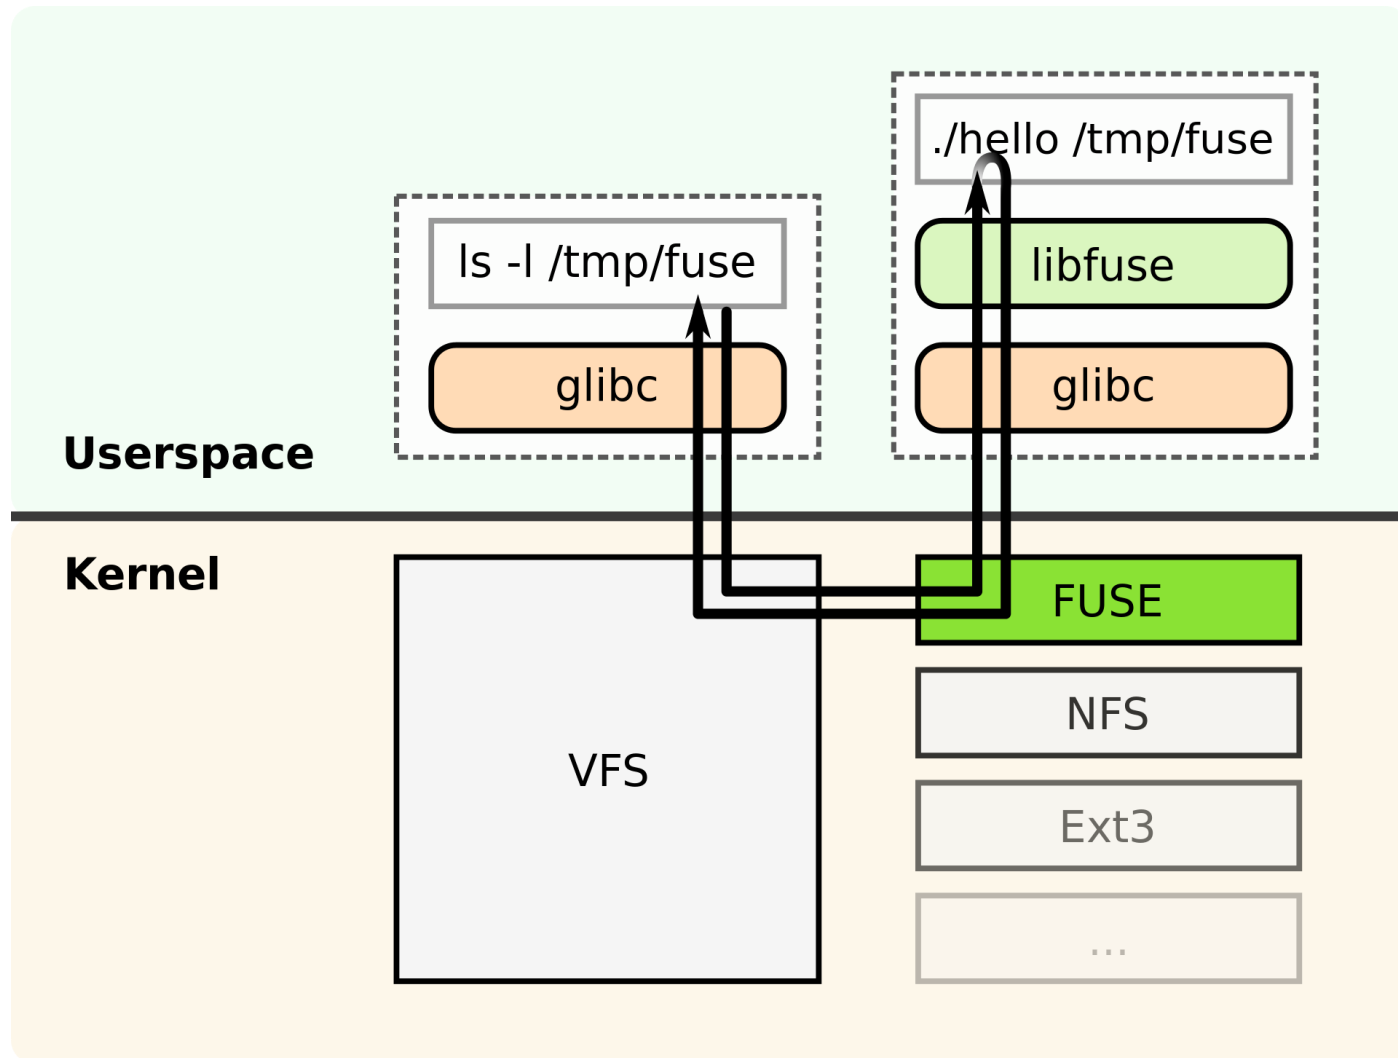
The Linux VFS

- Virtual File System (VFS)
 - defines a set of basic file-system abstractions and the operations which are allowed on these abstractions
 - Supports multiple file systems

Object	Description	Operation
Superblock	specific file-system	read_inode, sync_fs
Dentry	directory entry, single component of a path	create, link
I-node	specific file	d_compare, d_delete
File	open file associated with a process	read, write

Figure 10-30. File-system abstractions supported by the VFS.

File System in Userspace (FUSE)



Data Redundancy

□ Definition:

if A and B are two pieces of data,
and **knowing A eliminates some or all values B could be,**
there is redundancy between A and B

□ RAID examples:

- mirrored disk (complete redundancy)
- parity blocks (partial redundancy)

□ File system examples:

- **Superblock**: field contains total blocks in FS
- **Inodes**: field contains pointer to data block
- Is there redundancy between these two types of fields?
Why or why not?



File System Redundancy Example

Superblock: field contains total number of blocks in FS

DATA = N

Inode: field contains pointer to data block; possible DATA?

DATA in {0, 1, 2, ..., N - 1}

Pointers to block N or after are invalid!

Total-blocks field has **redundancy** with inode pointers

Consistency Examples

Assumptions:

Superblock: field contains total blocks in FS.

DATA = 1024

Inode: field contains pointer to data block.

DATA in {0, 1, 2, ..., 1023}

Scenario 1: Consistent or not?

Superblock: field contains total blocks in FS.

DATA = 1024

Inode: field contains pointer to data block.

DATA = 241

Consistent

Scenario 2: Consistent or not?

Superblock: field contains total blocks in FS.

DATA = 1024

node: field contains pointer to data block.

DATA = 2345

Inconsistent



Pros and CONs of Redundancy

Redundancy may improve:

- reliability

- RAID-5 parity
- Superblocks in FFS

- performance

- RAID-1 mirroring (reads)
- FFS group descriptor
- FFS bitmaps

Redundancy hurts:

- capacity

- consistency

- Redundancy implies certain combinations of values are illegal
- **Illegal combinations: inconsistency**



How to fix Inconsistencies?

□ Solution #1:

- **FSCK** = file system checker

FSCK can be pronounced "F-S-C-K", "F-S-check", "fizz-check", "F-sack", "fisk", "fishcake", "fizik", "F-sick", "F-sock", "F-sek", "feshk", etc.

□ Strategy:

- After crash, scan whole disk for contradictions and “fix” if needed
- Keep file system off-line until FSCK completes

□ For example, how to tell if data bitmap block is consistent?

Read **every** valid **inode+indirect** block


If pointer to data block, the corresponding bit should be 1;
else bit is 0

Fsck Checks

Hundreds of types of checks over different fields...


- ☐ Do superblocks match?
- ☐ Do directories contain “.” and “..”?
- ☐ Do number of dir entries equal **inode link counts**?
- ☐ Do different inodes ever point to **same block**?
- ☐ ...

Link Count (example 1)



How to fix to have consistent file system?

Link Count (example 1)





Link Count (example 2)

inode

link_count = 1

How to fix???

Link Count (example 2)

Dir Entry

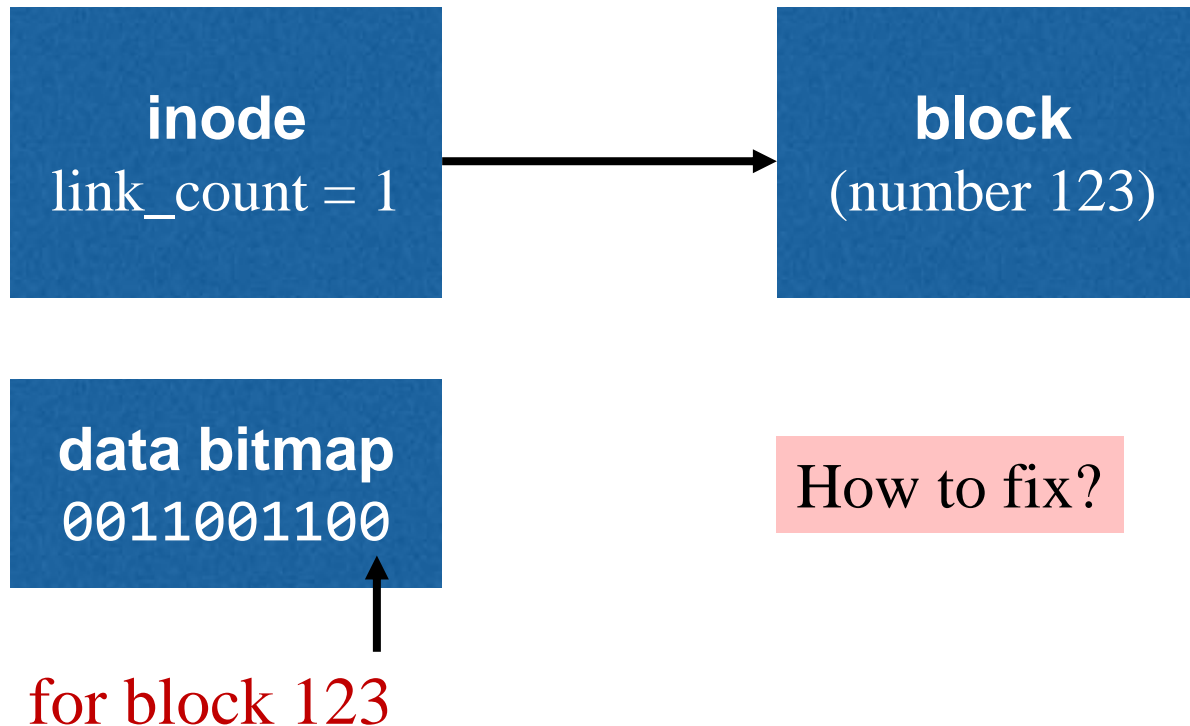
fix!

inode

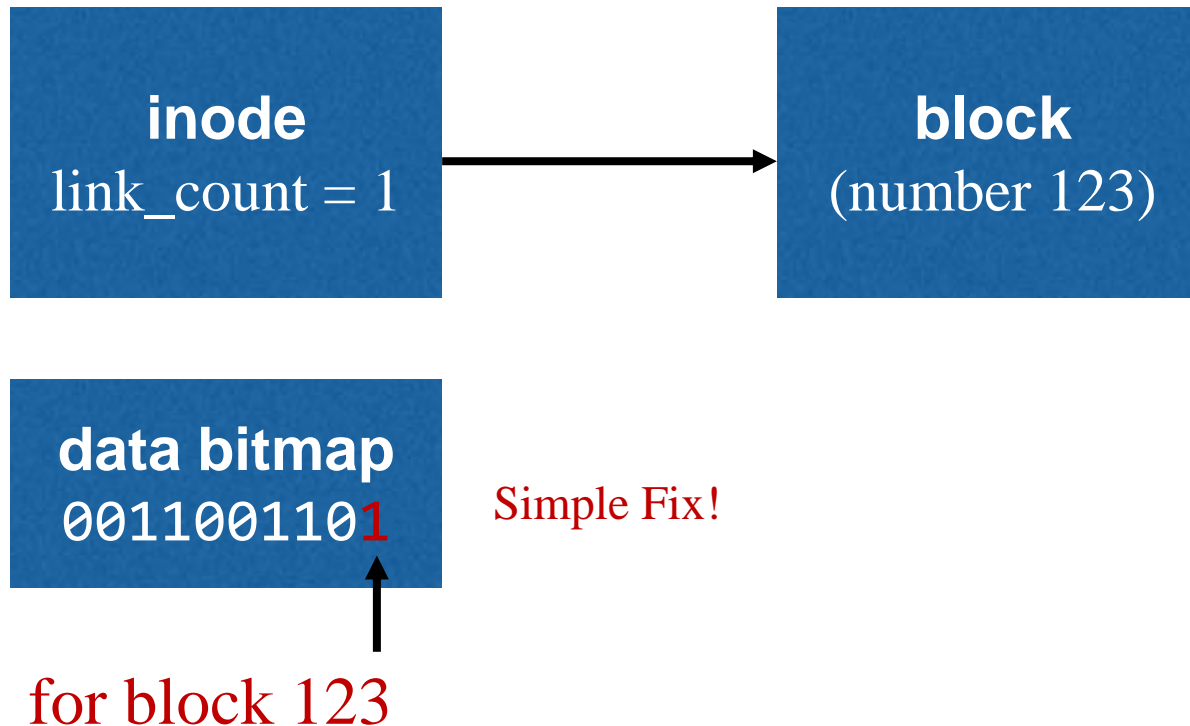
link_count = 1

```
ls -l /
total 150
drwxr-xr-x 401 18432 Dec 31 1969 afs/
drwxr-xr-x. 2 4096 Nov 3 09:42 bin/
drwxr-xr-x. 5 4096 Aug 1 14:21 boot/
dr-xr-xr-x. 13 4096 Nov 3 09:41 lib/
dr-xr-xr-x. 10 12288 Nov 3 09:41 lib64/
drwx-----. 2 16384 Aug 1 10:57 lost+found/
...
```

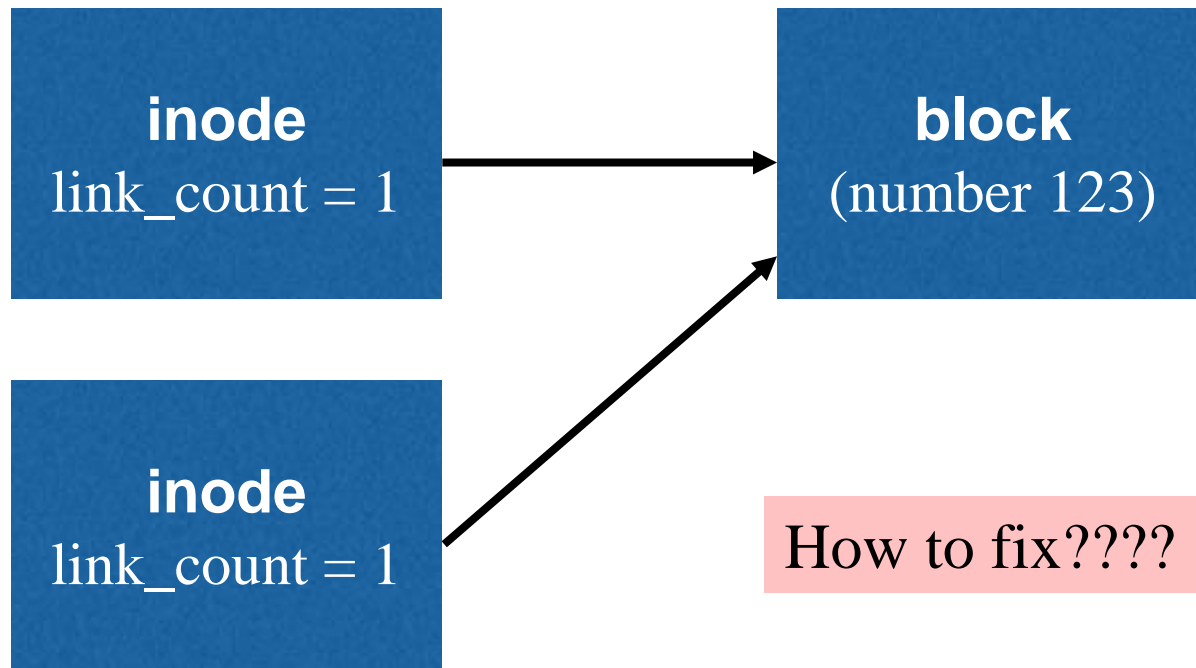

Data Bitmap




Data Bitmap




Duplicate Pointers



Duplicate Pointers

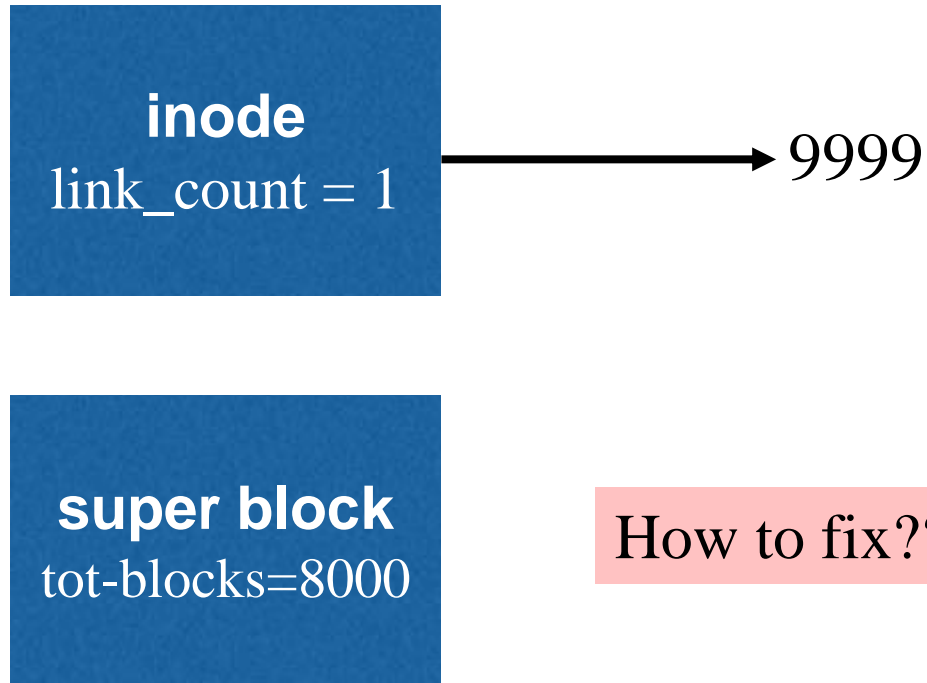


Duplicate Pointers



But is this correct?

Bad Pointer



Bad Pointer

inode
link_count = 1

Simple fix! (But is this correct?)

super block
tot-blocks=8000




Problems with fsck

Problem 1:

- Not always obvious how to fix file system image
- Don't know “correct” state, just consistent one
- Easy way to get consistency: reformat disk!

Problem 2: fsck is very slow



Checking a 600GB disk takes ~70 minutes

ffsck: The Fast File System Checker

Ao Ma, EMC Corporation and University of Wisconsin—Madison; Chris Dragga, Andrea C. Arpaci-Dusseau, and Remzi H. Arpaci-Dusseau, University of Wisconsin—Madison



Consistency Solution #2: Journaling

□ Goals

- Ok to do some **recovery work** after crash, but not to read entire disk
- Don't move file system to just any consistent state, get **correct** state

□ Strategy


- Atomicity
 - operations in critical sections are not interrupted by operations on related critical sections
- Definition of atomicity for **concurrency**
- Definition of atomicity for **persistence**
 - collections of writes are not interrupted by crashes; either (all new) or (all old) data is visible

Journaling General Strategy

- ❑ Never delete ANY old data, until, ALL new data is safely on disk
- ❑ Ironically, adding redundancy to fix the problem caused by redundancy.

Fight Redundancy with Redundancy


Want to replace X with Y. Original:





Fight Redundancy with Redundancy

Want to replace X with Y. Original:



Good time to crash?
good time to crash



Fight Redundancy with Redundancy

Want to replace X with Y. Original:

DISK



Good time to crash?
bad time to crash



Fight Redundancy with Redundancy

Want to replace X with Y. Original:

DISK



Good time to crash?
good time to crash



Fight Redundancy with Redundancy

Want to replace X with Y.

DISK




Good time to crash?
good time to crash



Fight Redundancy with Redundancy

Want to replace X with Y. With journal:



good time to crash



Fight Redundancy with Redundancy

Want to replace X with Y. With journal:

DISK



good time to crash



Fight Redundancy with Redundancy

Want to replace X with Y. With journal:

DISK



good time to crash



Fight Redundancy with Redundancy

Want to replace X with Y. With journal:

DISK



good time to crash



Fight Redundancy with Redundancy

Want to replace X with Y. With journal:

DISK



good time to crash



Fight Redundancy with Redundancy

Want to replace X with Y. With journal:

DISK




good time to crash



Fight Redundancy with Redundancy

Want to replace X with Y. With journal:



With journaling, it's always
a good time to crash!

Journaling is Widely Used

File System	Feature for block allocation
ext3 *	ext2 with journaling, Block Group is imported from FFS.
ext4 *	Successor of ext3, extent allocation , delayed allocation
JFS *	Dynamic i-node allocation, extent allocation .
XFS *	Variable block size, extent allocation .
ReiserFS (v3) *	Block sub-allocation(Tail packing)
Nilfs	stackable(log structured) FS
Btrfs	copy-on-write, extent allocation .
FAT32	FS for Windows, File allocation table. No journaling.
NTFS	FS for Windows NT, extent allocation . Linux uses NTFS-3G driver.

“*” indicates bootable FS.

All file systems except FAT32, have same function of journaling.

FAT

- File Allocation Table (FAT)
 - the entire table must be in memory

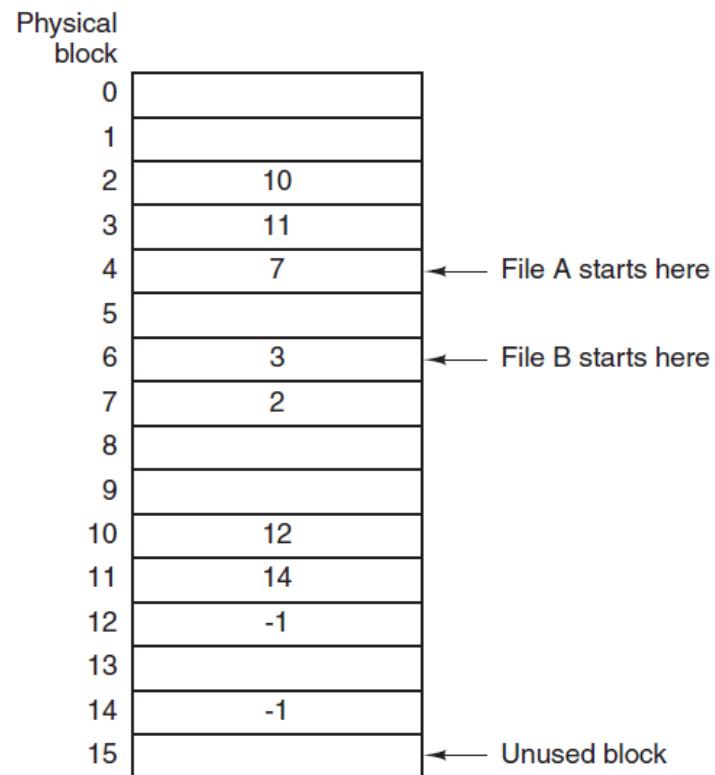


Figure 4-12. Linked-list allocation using a file-allocation table in main memory.

FAT: Used in MS-DOS and Windows

□ Directory entry structure




Figure 4-30. The MS-DOS directory entry.

Different FAT File Systems

- FAT-12, FAT-16, FAT-32
 - Uses different number of bits to address the blocks

Block size	FAT-12	FAT-16	FAT-32
0.5 KB	2 MB		
1 KB	4 MB		
2 KB	8 MB	128 MB	
4 KB	16 MB	256 MB	1 TB
8 KB		512 MB	2 TB
16 KB		1024 MB	2 TB
32 KB		2048 MB	2 TB

Figure 4-31. Maximum partition size for different block sizes. The empty boxes represent forbidden combinations.

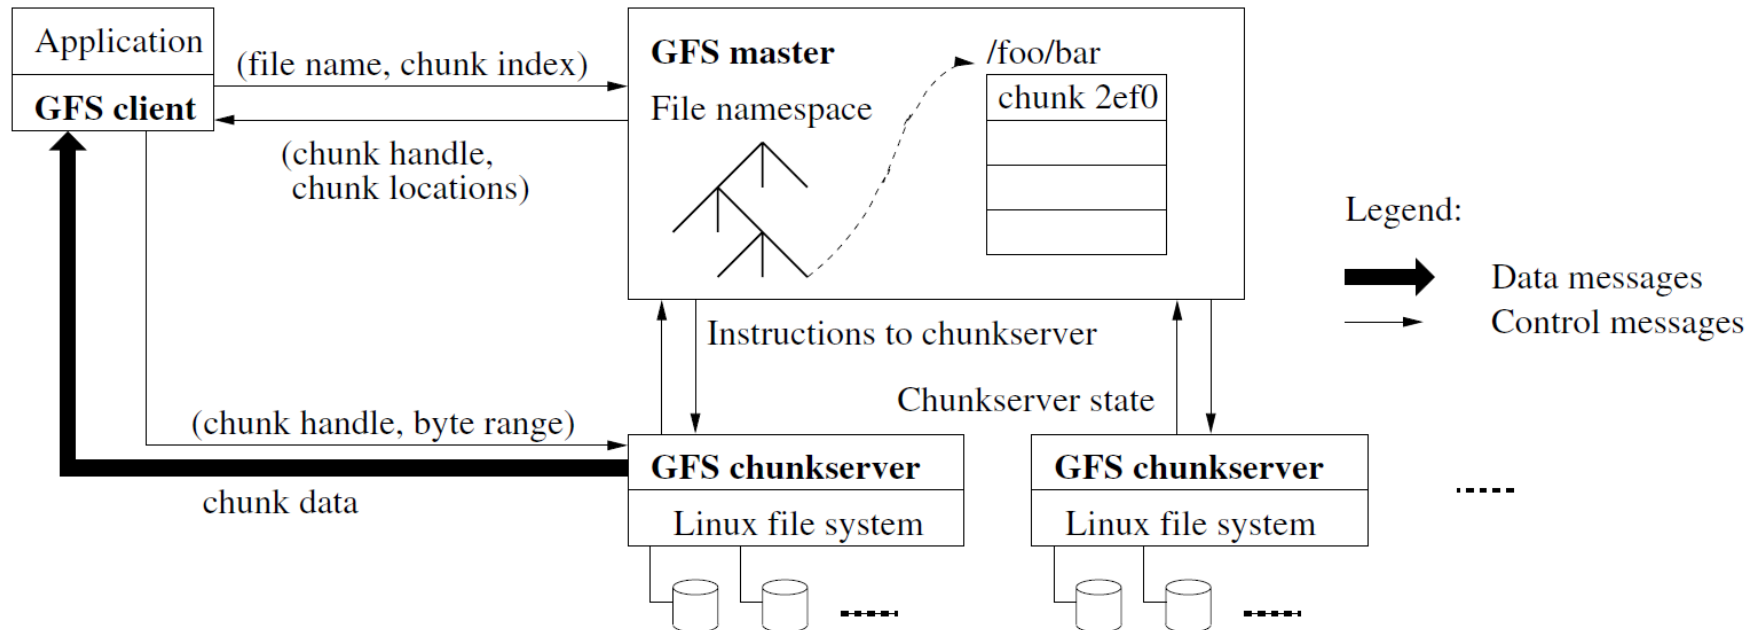
Google File System (GFS)

Motivation


- Google workload characteristics
 - huge files (GBs); usually read in their entirety
 - almost all writes are appends
 - concurrent appends common
 - high throughput is valuable
 - low latency is not
- Computing environment:
 - 1000s of machines
 - Machines sometimes fail (both permanently and temporarily)

GFS Architecture


- chunks: 64MB each



GFS Performance



(a) Reads



(b) Writes

GFS Clusters

- Note: paper published in SOSP 2003

Cluster	A	B
Chunkservers	342	227
Available disk space	72 TB	180 TB
Used disk space	55 TB	155 TB
Number of Files	735 k	737 k
Number of Dead files	22 k	232 k
Number of Chunks	992 k	1550 k
Metadata at chunkservers	13 GB	21 GB
Metadata at master	48 MB	60 MB

GFS Performance

Cluster	A	B
Read rate (last minute)	583 MB/s	380 MB/s
Read rate (last hour)	562 MB/s	384 MB/s
Read rate (since restart)	589 MB/s	49 MB/s
Write rate (last minute)	1 MB/s	101 MB/s
Write rate (last hour)	2 MB/s	117 MB/s
Write rate (since restart)	25 MB/s	13 MB/s
Master ops (last minute)	325 Ops/s	533 Ops/s
Master ops (last hour)	381 Ops/s	518 Ops/s
Master ops (since restart)	202 Ops/s	347 Ops/s

Table 3: Performance Metrics for Two GFS Clusters

Summary

- File system is evolving rapidly
- New file systems for new computing environments
 - FAST: USENIX Conference on File and Storage Technologies
- Next lecture
 - I/O